

# Shallow Neural Networks

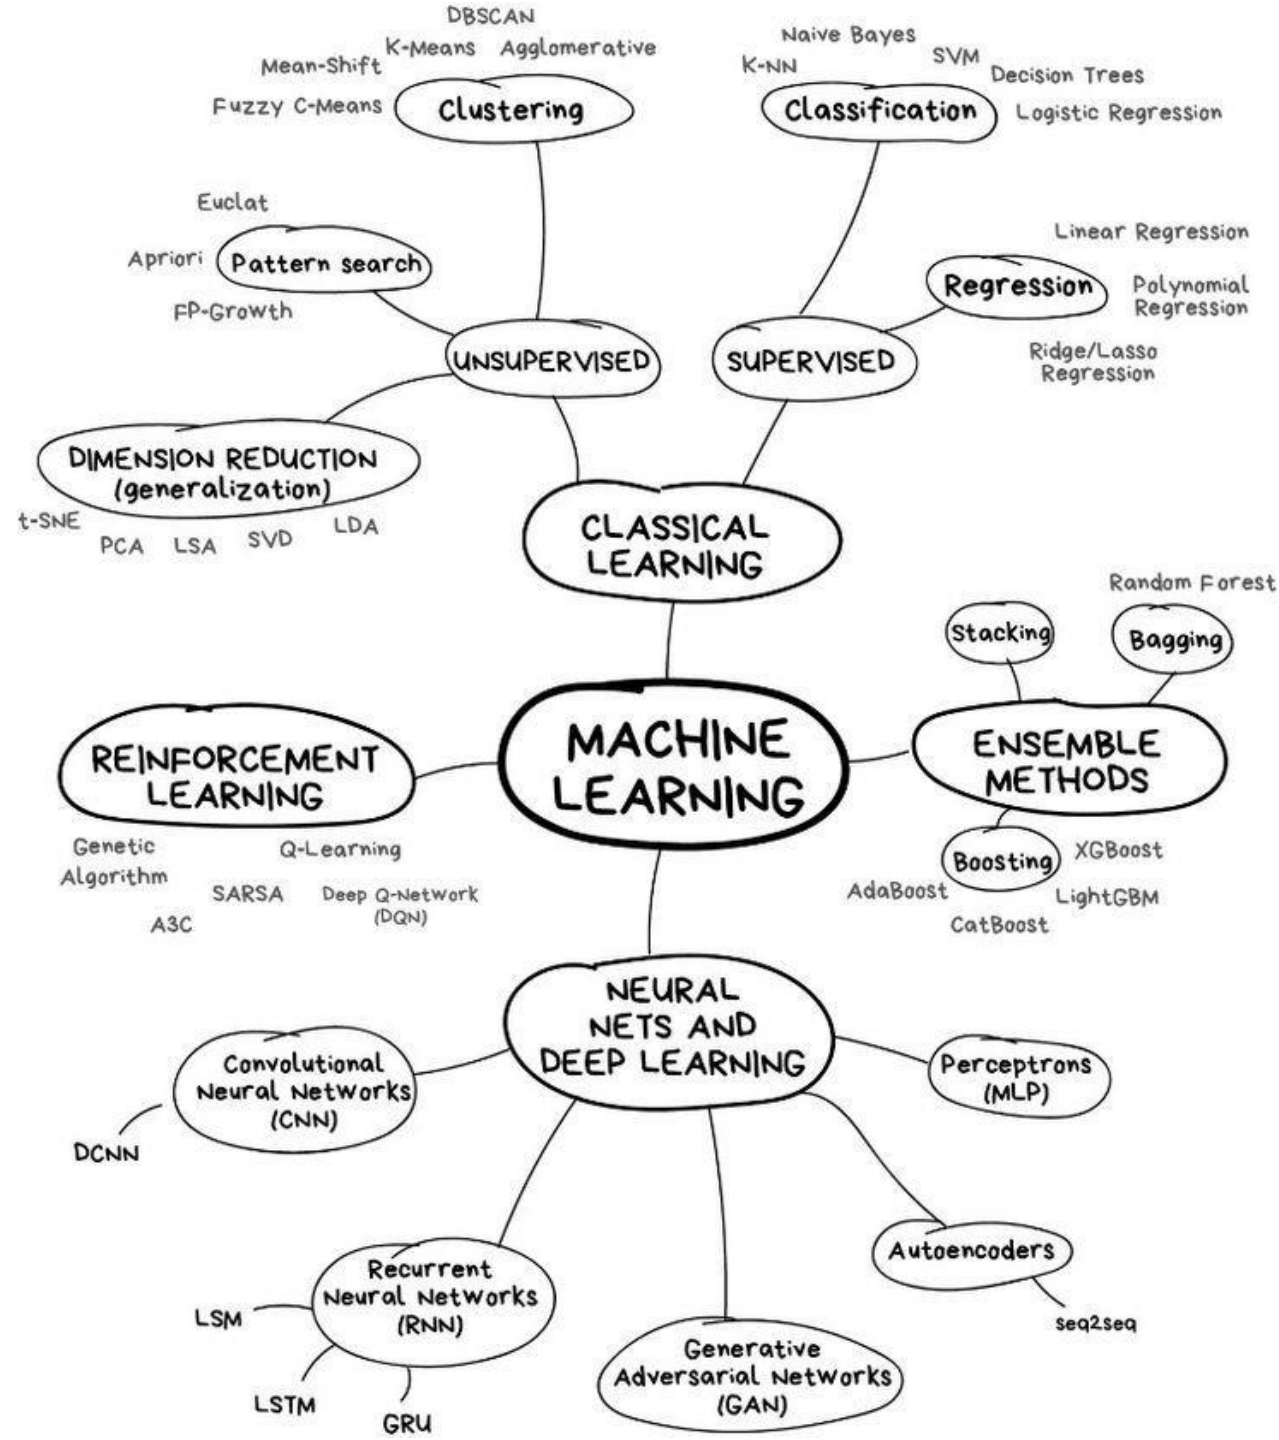


한양대학교 ERICA  
소프트웨어융합대학  
COLLEGE OF COMPUTING

인공지능학과  
Department of  
Artificial Intelligence



정 우 환 (whjung@hanyang.ac.kr)

Fall 2021




# Logistic Regression


- Output:  $\hat{y} = \sigma(\mathbf{w}^\top \mathbf{x} + b)$  where  $\sigma(z) = \frac{1}{1+e^{-z}}$
- Loss:  $L(\hat{y}, y) = -y \log \hat{y} - (1 - y) \log(1 - \hat{y})$




# Decision Boundaries




kNN, k=5




kNN, k=15




Logistic Regression  
simple



Logistic Regression  
basic polynomials




Decision tree




SVM

Random forest



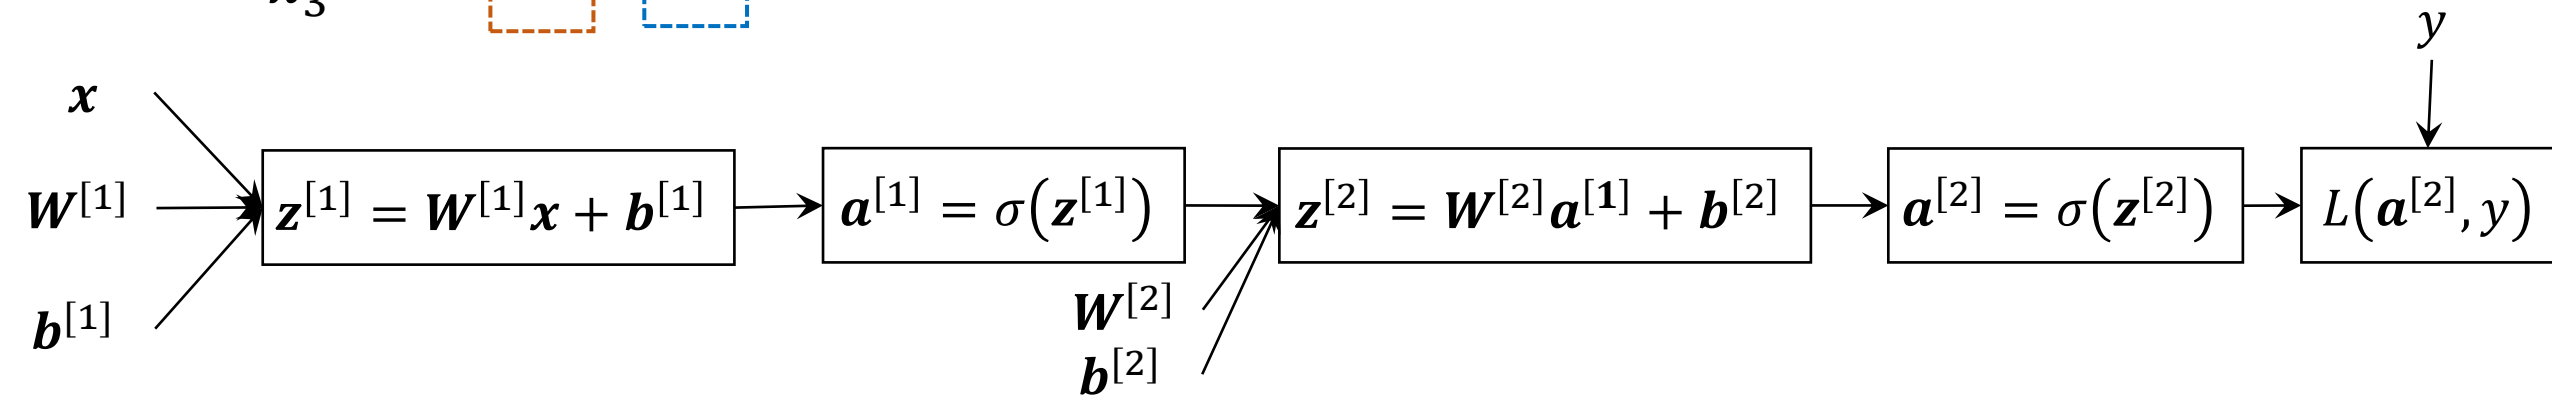


SVM

(Simple) XOR problem: linearly separable?



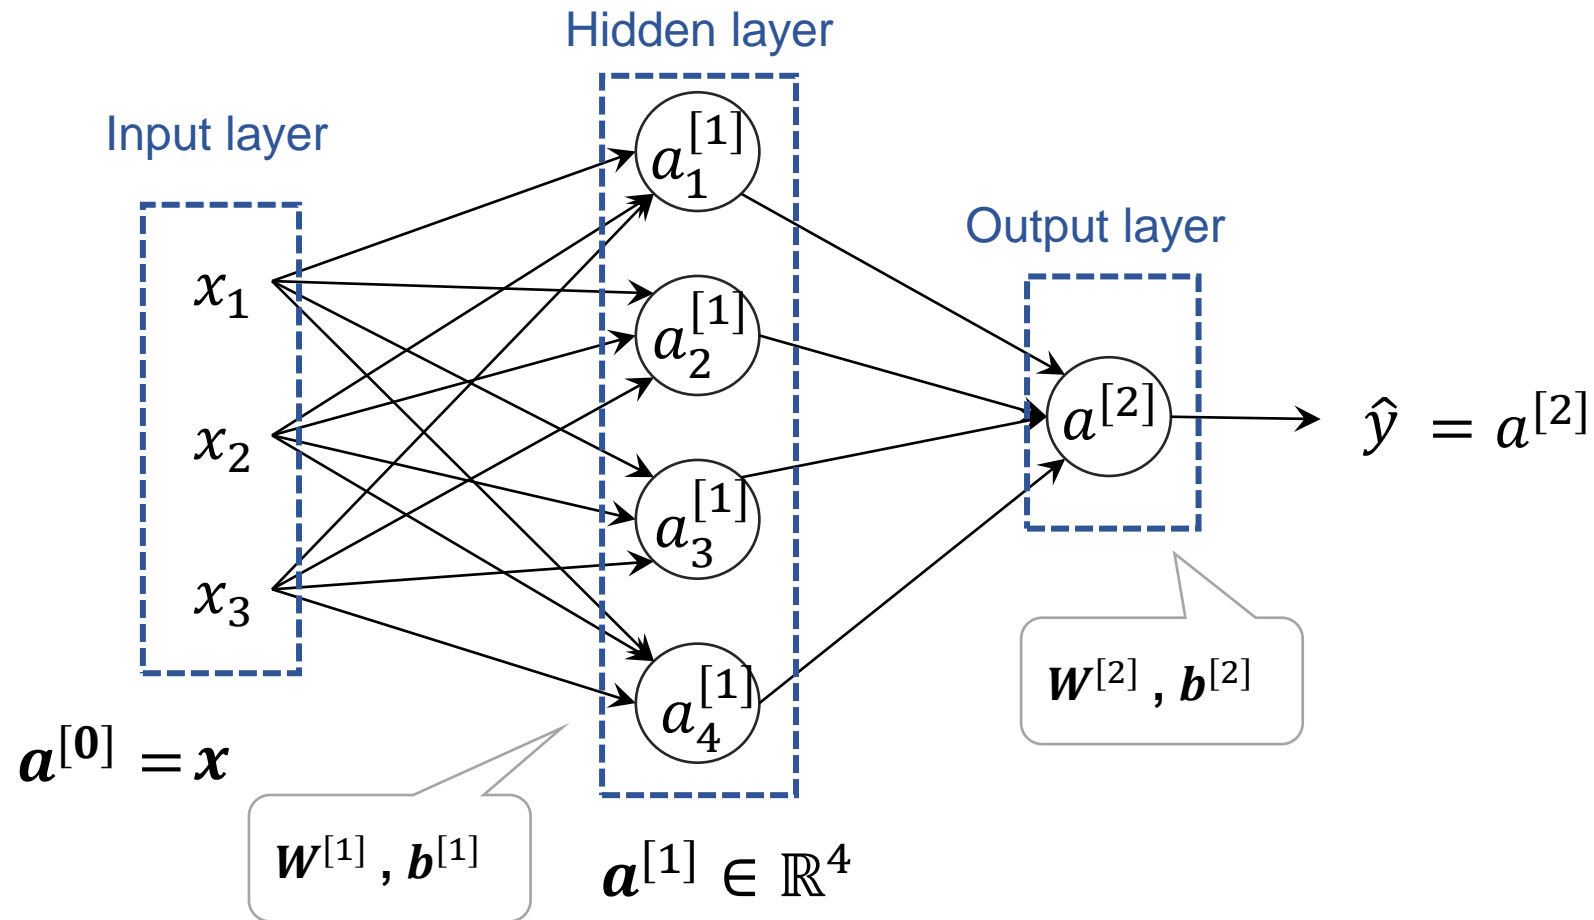
**Solution: make it more complicated**

# What is a Neural Network?




# Neural Network Representation

- 2-layer neural network




# Neural Network Representation




$$z = w^T x + b$$

$$a = \sigma(z)$$





# Neural Network Representation



$$z_1^{[1]} = \mathbf{w}_1^{[1]T} \mathbf{x} + b_1^{[1]} \quad a_1^{[1]} = \sigma(z_1^{[1]})$$

$$z_2^{[1]} = \mathbf{w}_2^{[1]T} \mathbf{x} + b_2^{[1]} \quad a_2^{[1]} = \sigma(z_2^{[1]})$$

$$z_3^{[1]} = \mathbf{w}_3^{[1]T} \mathbf{x} + b_3^{[1]} \quad a_3^{[1]} = \sigma(z_3^{[1]})$$

$$z_4^{[1]} = \mathbf{w}_4^{[1]T} \mathbf{x} + b_4^{[1]} \quad a_4^{[1]} = \sigma(z_4^{[1]})$$

$$z^{[2]} = \mathbf{w}^{[2]T} \mathbf{a}^{[1]} + b^{[2]} \quad a^{[2]} = \sigma(z^{[2]})$$

# Neural Network Representation: Vectorization

Hidden layer

$$z_i^{[1]} = \mathbf{W}_i^{[1]} \mathbf{x} + b_i^{[1]} \quad a_i^{[1]} = \sigma(z_i^{[1]}) \quad \text{For } 1 \leq i \leq h \text{ where } h: \# \text{ of hidden nodes}$$


Output layer

$$z^{[2]} = \mathbf{w}^{[2]T} \mathbf{a}^{[1]} + b^{[2]} \quad a^{[2]} = \sigma(z^{[2]})$$

Let  $\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \dots \\ x_m \end{bmatrix}$   $\mathbf{w}_i^{[1]} \in \mathbb{R}^m$

$$\mathbf{z}^{[1]} = \begin{bmatrix} z_1^{[1]} \\ z_2^{[1]} \\ \dots \\ z_h^{[1]} \end{bmatrix} \quad \mathbf{b}^{[1]} = \begin{bmatrix} b_1^{[1]} \\ b_2^{[1]} \\ \dots \\ b_h^{[1]} \end{bmatrix} \quad \mathbf{W}^{[1]} = \begin{bmatrix} \mathbf{w}_1^{[1]T} \\ \mathbf{w}_2^{[1]T} \\ \dots \\ \mathbf{w}_h^{[1]T} \end{bmatrix} = \begin{bmatrix} w_{1,1}^{[1]} & \dots & w_{1,m}^{[1]} \\ \vdots & \ddots & \vdots \\ w_{h,1}^{[1]} & \dots & w_{h,m}^{[1]} \end{bmatrix} \quad \mathbf{a}^{[1]} = \begin{bmatrix} a_1^{[1]} \\ a_2^{[1]} \\ \dots \\ a_h^{[1]} \end{bmatrix} \quad \mathbf{w}^{[2]} = \begin{bmatrix} w_1^{[2]} \\ w_2^{[2]} \\ \dots \\ w_h^{[2]} \end{bmatrix}$$

# Neural Network with a Hidden Layer: Almost Done!



Input:

$$\mathbf{x} \in \mathbb{R}^m$$

Parameters:

$$\mathbf{W}^{[1]} \in \mathbb{R}^{h \times m}$$

$$\mathbf{b}^{[1]} \in \mathbb{R}^h$$

$$\mathbf{w}^{[2]} \in \mathbb{R}^h$$

$$b^{[2]} \in \mathbb{R}$$

Forward pass:

$$\mathbf{z}^{[1]} = \mathbf{W}^{[1]} \mathbf{x} + \mathbf{b}^{[1]}$$


$$\mathbf{a}^{[1]} = \sigma(\mathbf{z}^{[1]})$$

$$z^{[2]} = \mathbf{w}^{[2]T} \mathbf{a}^{[1]} + b^{[2]}$$

$$\hat{y} = a^{[2]} = \sigma(z^{[2]})$$

# Activation Functions

# Neural Network with a Hidden Layer: Almost Done!



Input:

$$\mathbf{x} \in \mathbb{R}^m$$

Parameters:

$$\begin{aligned} \mathbf{W}^{[1]} &\in \mathbb{R}^{h \times m} & \mathbf{w}^{[2]} &\in \mathbb{R}^h \\ \mathbf{b}^{[1]} &\in \mathbb{R}^h & b^{[2]} &\in \mathbb{R} \end{aligned}$$

Forward pass:

$$\mathbf{z}^{[1]} = \mathbf{W}^{[1]} \mathbf{x} + \mathbf{b}^{[1]}$$

$$\mathbf{a}^{[1]} = g(\mathbf{z}^{[1]}) \text{ where } g(.) \text{ is an activation function}$$

$$\mathbf{z}^{[2]} = \mathbf{w}^{[2]T} \mathbf{a}^{[1]} + b^{[2]}$$

$$\hat{y} = a^{[2]} = \sigma(\mathbf{z}^{[2]})$$

# Why we need to use non-linear activation functions?



$$\begin{aligned} z^{[2]} &= W^{[2]}(W^{[1]}x + b^{[1]}) + b^{[2]} \\ &= W^{[2]}W^{[1]}x + (W^{[2]}b^{[1]} + b^{[2]}) \\ &= W'x + b' \end{aligned}$$

$$\text{Where } W' = W^{[2]}W^{[1]} \text{ and } b' = W^{[2]}b^{[1]} + b^{[2]}$$

Composition of linear functions => linear function

# Activation functions

There are so many activation functions ..

We'll cover some important and currently widely used activation functions

Sigmoid

tanh

ReLU

LeakyReLU


Name	Plot	Function, $f(x)$	Derivative of $f$ , $f'(x)$
Identity		$x$	1
Binary step		$\begin{cases} 0 & \text{if } x < 0 \\ 1 & \text{if } x \geq 0 \end{cases}$	$\begin{cases} 0 & \text{if } x \neq 0 \\ \text{undefined} & \text{if } x = 0 \end{cases}$
Logistic, sigmoid, or soft step		$\sigma(x) = \frac{1}{1 + e^{-x}}$ [1]	$f(x)(1 - f(x))$
tanh		$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$	$1 - f(x)^2$
Rectified linear unit (ReLU) [11]		$\begin{cases} 0 & \text{if } x \leq 0 \\ x & \text{if } x > 0 \end{cases}$ $= \max\{0, x\} = x \mathbf{1}_{x>0}$	$\begin{cases} 0 & \text{if } x < 0 \\ 1 & \text{if } x > 0 \\ \text{undefined} & \text{if } x = 0 \end{cases}$
Gaussian Error Linear Unit (GELU) [6]		$\frac{1}{2}x \left(1 + \operatorname{erf}\left(\frac{x}{\sqrt{2}}\right)\right)$ $= x\Phi(x)$	$\Phi(x) + x\phi(x)$
Softplus [12]		$\ln(1 + e^x)$	$\frac{1}{1 + e^{-x}}$
Exponential linear unit (ELU) [13]		$\begin{cases} \alpha(e^x - 1) & \text{if } x \leq 0 \\ x & \text{if } x > 0 \end{cases}$ with parameter $\alpha$	$\begin{cases} \alpha e^x & \text{if } x < 0 \\ 1 & \text{if } x > 0 \\ 1 & \text{if } x = 0 \text{ and } \alpha = 1 \end{cases}$
Square linear unit (SQU) [14]		$\begin{cases} x & \text{if } x > 0.0 \\ \alpha(x + \frac{x^2}{4}) & \text{if } -2.0 \leq x \leq 0 \\ -\alpha & \text{if } x < -2.0 \end{cases}$	$\begin{cases} 1 & \text{if } x > 0.0 \\ 1 + \frac{x}{2} & \text{if } -2.0 \leq x \leq 0 \\ 0 & \text{if } x < -2.0 \end{cases}$
Scaled exponential linear unit (SELU) [15]		$\lambda \begin{cases} \alpha(e^x - 1) & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$ with parameters $\lambda = 1.0507$ and $\alpha = 1.67326$	$\lambda \begin{cases} \alpha e^x & \text{if } x < 0 \\ 1 & \text{if } x \geq 0 \end{cases}$
Leaky rectified linear unit (Leaky ReLU) [16]		$\begin{cases} 0.01x & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$	$\begin{cases} 0.01 & \text{if } x < 0 \\ 1 & \text{if } x \geq 0 \end{cases}$
Parametric rectified linear unit (PReLU) [17]		$\begin{cases} \alpha x & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$ with parameter $\alpha$	$\begin{cases} \alpha & \text{if } x < 0 \\ 1 & \text{if } x \geq 0 \end{cases}$
ElliotSig, [18][19] softsign [20][21]		$\frac{x}{1 +  x }$	$\frac{1}{(1 +  x )^2}$

# Sigmoid

- $\sigma(x) = \frac{1}{1+e^{-x}}$
- Range: (0,1)
- Derivative

$$\frac{d\sigma(x)}{dx} = \sigma(x)(1 - \sigma(x))$$

- $0 < \frac{d\sigma(x)}{dx} \leq 0.25$





# Hyperbolic tangent: tanh


- $\tanh x = \frac{e^x - e^{-x}}{e^x + e^{-x}}$   
 $= 2\sigma(2x) - 1$

- Range:  $(-1, 1)$

- Derivative

$$\frac{d \tanh x}{dx} = 1 - \tanh^2 x$$

- $0 < \frac{d \tanh x}{dx} \leq 1$



# Vanishing gradient

“The term vanishing gradient refers to the fact that in a feedforward network (FFN) the backpropagated error signal typically decreases (or increases) exponentially as a function of the distance from the final layer” by Jason Brownlee

<https://machinelearningmastery.com/how-to-fix-vanishing-gradients-using-the-rectified-linear-activation-function/>



- Chain rule (for a Deep NN)

$$\frac{\partial L}{\partial z^{[n]}} \frac{\partial z^{[n]}}{\partial z^{[n-1]}} \cdots \frac{\partial z^{[3]}}{\partial z^{[2]}} \frac{\partial z^{[2]}}{\partial z^{[1]}} \frac{\partial z^{[1]}}{\partial w}$$

# Rectified linear unit: ReLU

- $f(x) = \max\{0, x\}$
- Range:  $[0, \infty)$
- Derivative

$$\frac{df(x)}{dx} = \begin{cases} 0 & \text{if } x < 0 \\ 1 & \text{if } x > 0 \\ \text{undefined} & \text{if } x = 0 \end{cases}$$




# Leaky ReLU


- $f(x) = \begin{cases} ax & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$ 
  - $a \ll 1$  (e.g,  $a = 0.01$ )
- Range:  $(-\infty, \infty)$
- Derivative

$$\frac{df(x)}{dx} = \begin{cases} a & \text{if } x < 0 \\ 1 & \text{if } x > 0 \\ \text{undefined} & \text{if } x = 0 \end{cases}$$

ReLU



Leaky ReLU




# Activation functions

- There is no rule but ... in many cases
- Output layer
  - Sigmoid
- Hidden layer
  - tanh, ReLU, LeakyReLU

My personal opinion:  
If the performance of two things are similar,  
**Use the simpler one! (Use ReLU!)**

# Neural Network with a Hidden Layer: Done!



$$\mathbf{z}^{[1]} = \mathbf{W}^{[1]}\mathbf{x} + \mathbf{b}^{[1]}$$

$$\mathbf{a}^{[1]} = g(\mathbf{z}^{[1]}) \text{ where } g(.) \text{ is an activation function}$$

$$\mathbf{z}^{[2]} = \mathbf{w}^{[2]T} \mathbf{a}^{[1]} + b^{[2]}$$


$$\hat{y} = a^{[2]} = \sigma(\mathbf{z}^{[2]})$$

You may use tanh, ReLU, Leaky ReLU for  $g$

# Gradient Descent

Shallow Neural Network with tanh

# Shallow Neural Network



$$\mathbf{z}^{[1]} = \mathbf{W}^{[1]}\mathbf{x} + \mathbf{b}^{[1]}$$

$$\mathbf{a}^{[1]} = \tanh \mathbf{z}^{[1]}$$

$$z^{[2]} = \mathbf{w}^{[2]T} \mathbf{a}^{[1]} + b^{[2]}$$

$$\hat{y} = a^{[2]} = \sigma(z^{[2]})$$

Parameters:  $\boldsymbol{\theta} = \{\mathbf{W}^{[1]}, \mathbf{b}^{[1]}, \mathbf{w}^{[2]}, b^{[2]}\}$

Loss  $L(\hat{y}, y)$

Cost  $J(\boldsymbol{\theta})$

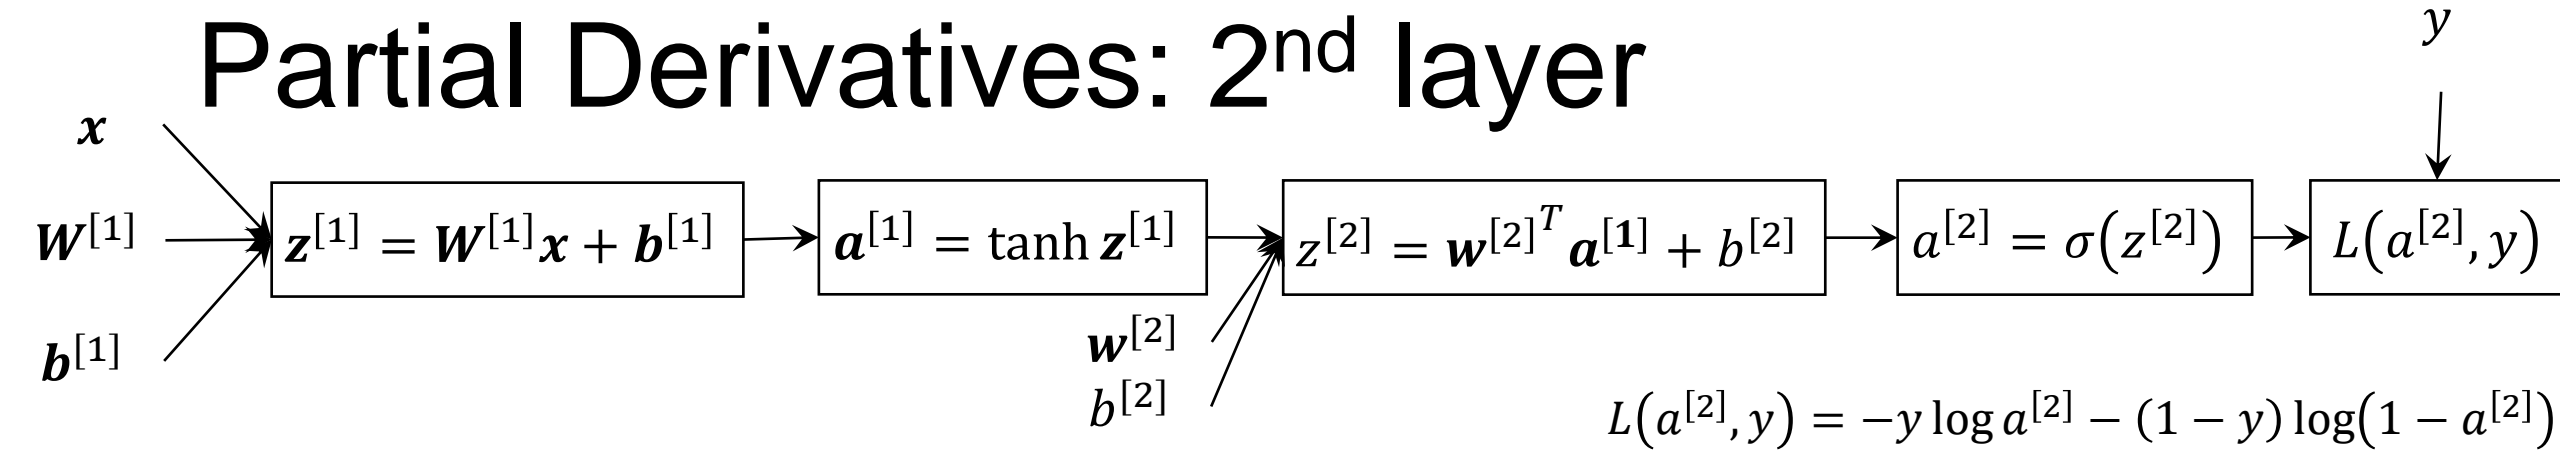
Gradient descent

$$\boldsymbol{\theta} := \boldsymbol{\theta} - \eta \nabla J = \boldsymbol{\theta} - \eta \frac{\nabla L}{m}$$



Parameters  $W^{[1]}, b^{[1]}, w^{[2]}, b^{[2]}$

# Partial Derivatives: 2<sup>nd</sup> layer



$$L(a^{[2]}, y) = -y \log a^{[2]} - (1 - y) \log(1 - a^{[2]})$$

$$\bullet \frac{\partial L(a^{[2]}, y)}{\partial b^{[2]}} = \frac{\partial L(a^{[2]}, y)}{\partial a^{[2]}} \frac{\partial a^{[2]}}{\partial z^{[2]}} \frac{\partial z^{[2]}}{\partial b^{[2]}}$$

$$\frac{d\sigma(x)}{dx} = \sigma(x)(1 - \sigma(x))$$

$$= \left( \frac{-y}{a^{[2]}} + \frac{1 - y}{1 - a^{[2]}} \right) \sigma(z^{[2]}) (1 - \sigma(z^{[2]}))$$

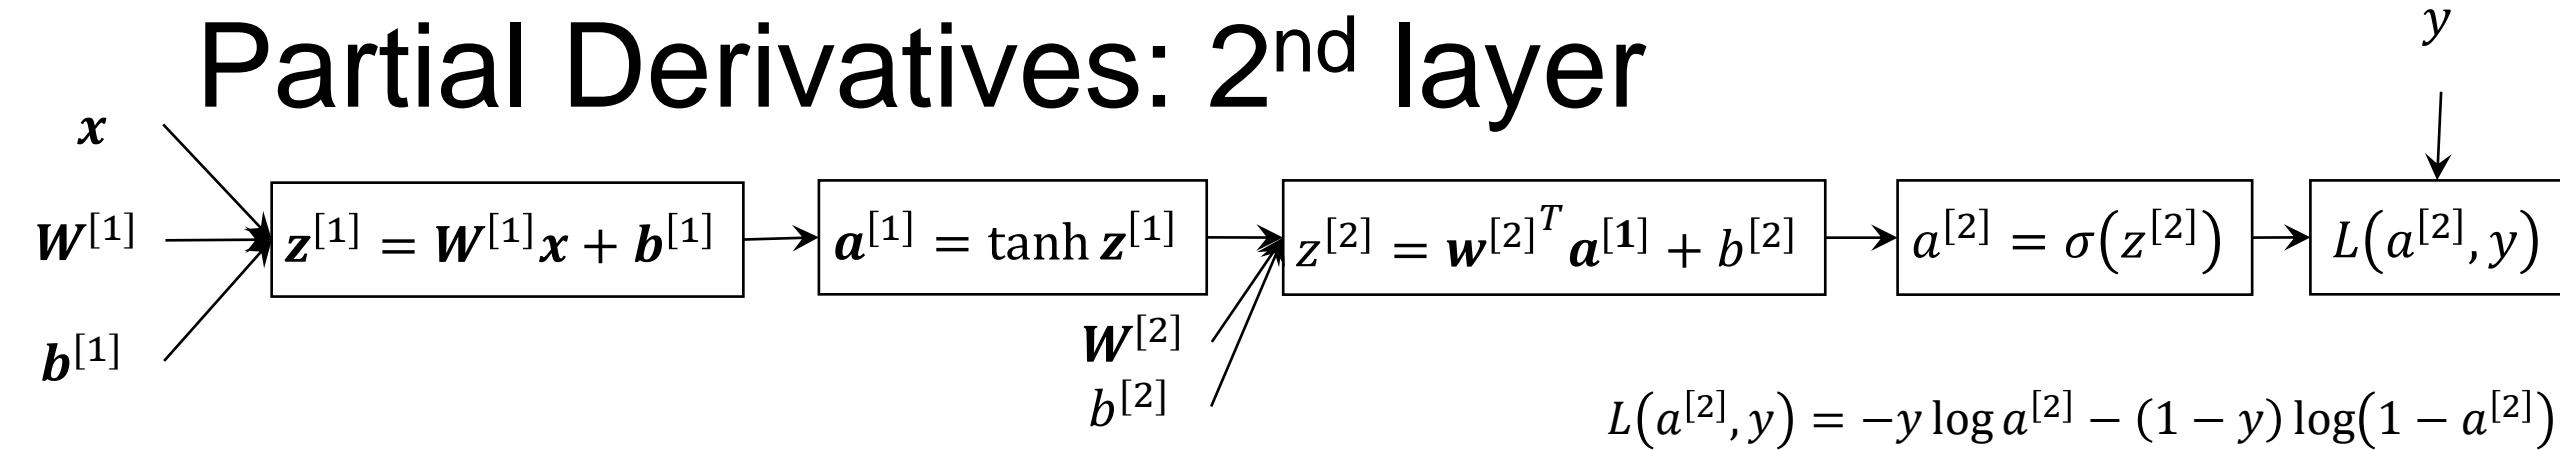
$$= \left( \frac{-y}{a^{[2]}} + \frac{1 - y}{1 - a^{[2]}} \right) a^{[2]} (1 - a^{[2]})$$

$$= -y(1 - a^{[2]}) + a^{[2]}(1 - y) = a^{[2]} - y$$

Do you remember?

Parameters  $\mathbf{W}^{[1]}, \mathbf{b}^{[1]}, \mathbf{w}^{[2]}, b^{[2]}$

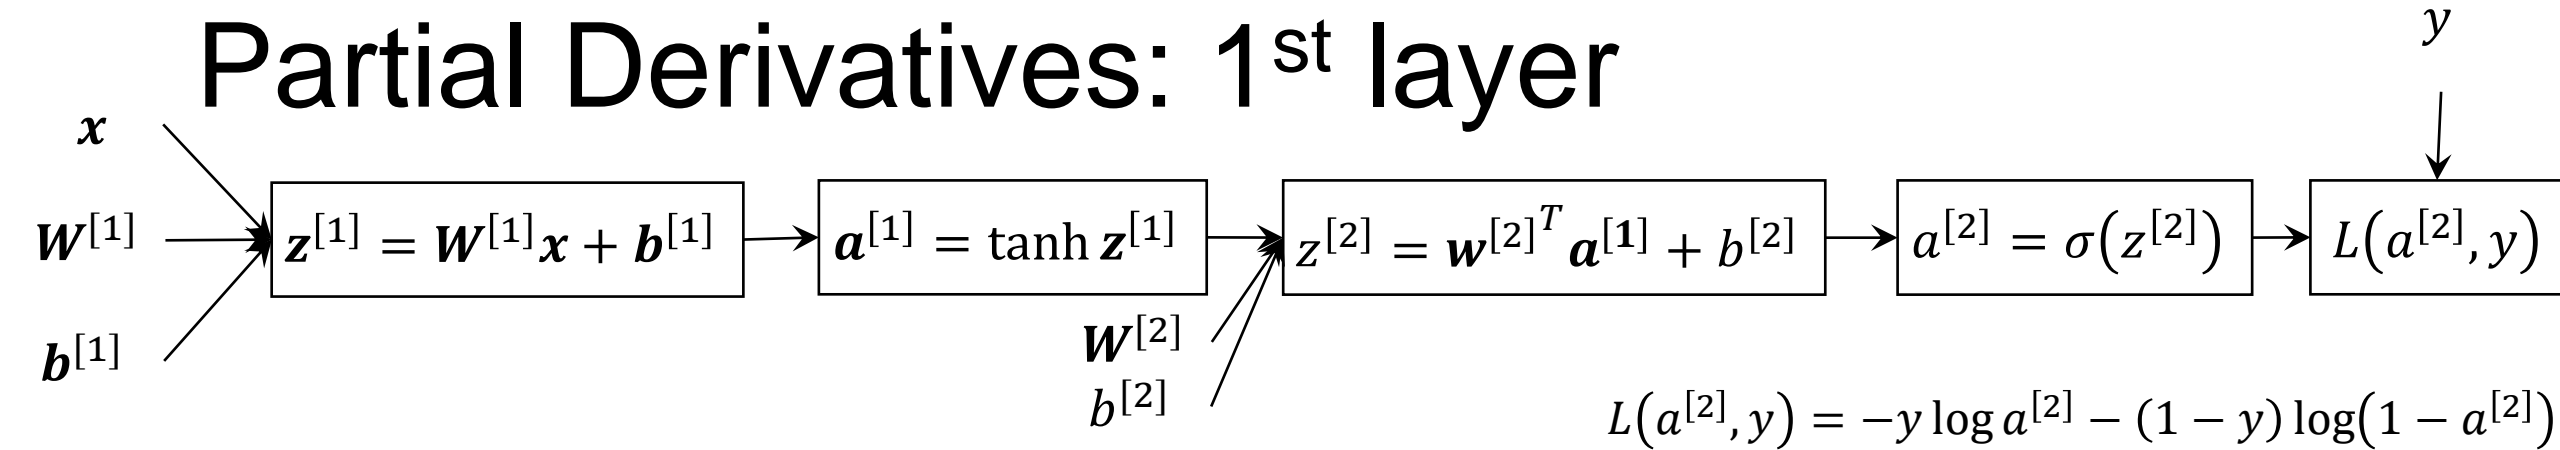
# Partial Derivatives: 2<sup>nd</sup> layer



$$\begin{aligned}
 \bullet \frac{\partial L(a^{[2]}, y)}{\partial w_i^{[2]}} &= \frac{\partial L(a^{[2]}, y)}{\partial a^{[2]}} \frac{\partial a^{[2]}}{\partial z^{[2]}} \frac{\partial z^{[2]}}{\partial w_i^{[2]}} \\
 &= (a^{[2]} - y) \frac{\partial z^{[2]}}{\partial w_i^{[2]}} \\
 &= (a^{[2]} - y) a_i^{[1]}
 \end{aligned}$$

Parameters  $\mathbf{W}^{[1]}, \mathbf{b}^{[1]}, \mathbf{w}^{[2]}, b^{[2]}$

# Partial Derivatives: 1<sup>st</sup> layer

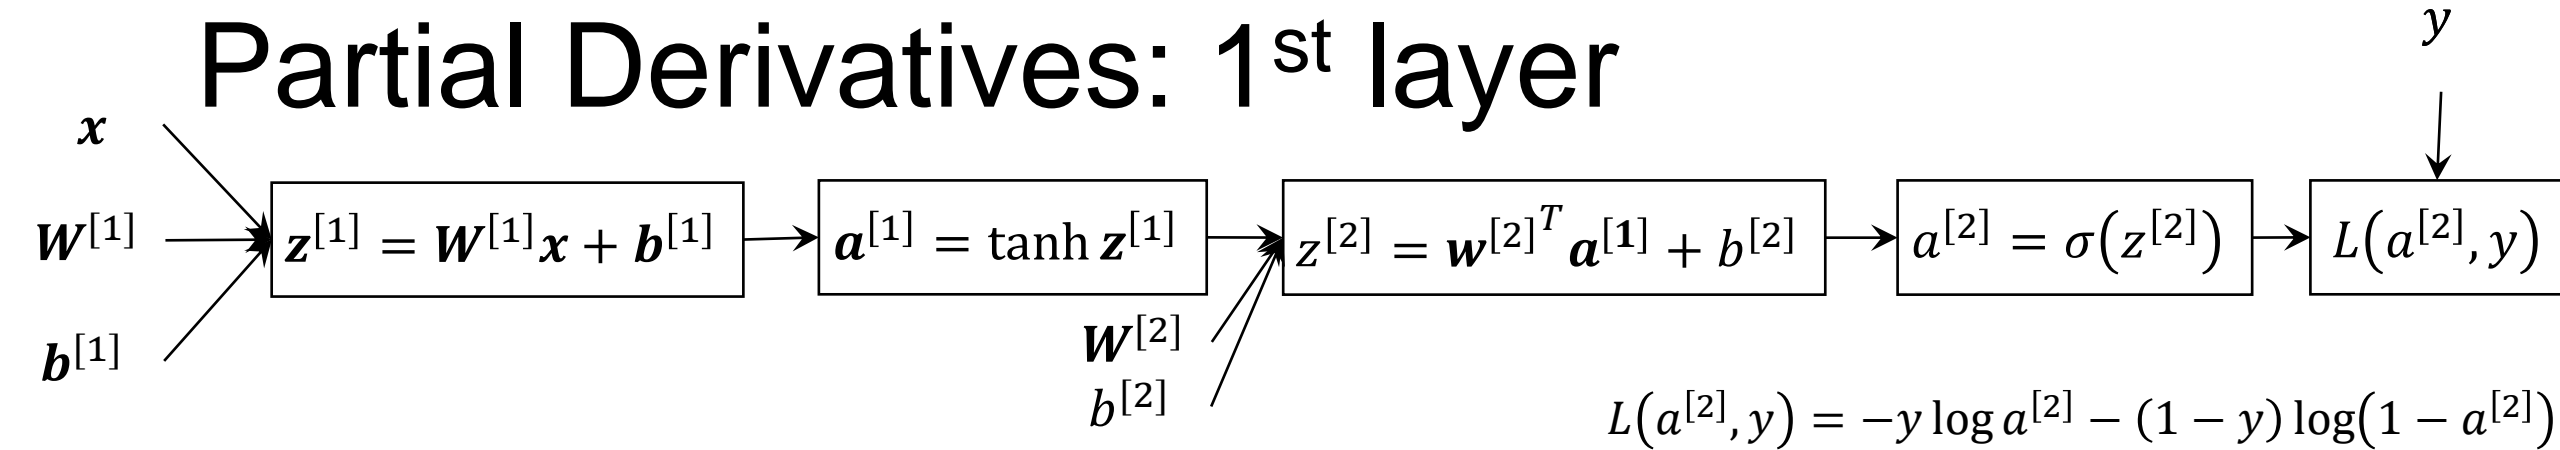


$$\begin{aligned}
 \bullet \frac{\partial L(a^{[2]}, y)}{\partial b_i^{[1]}} &= \frac{\partial L(a^{[2]}, y)}{\partial a^{[2]}} \frac{\partial a^{[2]}}{\partial z^{[2]}} \frac{\partial z^{[2]}}{\partial a_i^{[1]}} \frac{\partial a_i^{[1]}}{\partial z_i^{[1]}} \frac{\partial z_i^{[1]}}{\partial b_i^{[1]}} \\
 &= (a^{[2]} - y) \frac{\partial z^{[2]}}{\partial a_i^{[1]}} \frac{\partial a_i^{[1]}}{\partial z_i^{[1]}} \frac{\partial z_i^{[1]}}{\partial b_i^{[1]}} \\
 &= (a^{[2]} - y) w_i^{[2]} (1 - \tanh^2 z_i^{[1]}) \cdot 1 \\
 &= (a^{[2]} - y) w_i^{[2]} (1 - a_i^{[1]2})
 \end{aligned}$$

$$\frac{d \tanh x}{dx} = 1 - \tanh^2 x$$

Parameters  $\mathbf{W}^{[1]}, \mathbf{b}^{[1]}, \mathbf{w}^{[2]}, b^{[2]}$

# Partial Derivatives: 1<sup>st</sup> layer



$$L(a^{[2]}, y) = -y \log a^{[2]} - (1 - y) \log(1 - a^{[2]})$$

$$\begin{aligned} \bullet \frac{\partial L(a^{[2]}, y)}{\partial W_{ij}^{[1]}} &= \frac{\partial L(a^{[2]}, y)}{\partial a^{[2]}} \frac{\partial a^{[2]}}{\partial z^{[2]}} \frac{\partial z^{[2]}}{\partial a_i^{[1]}} \frac{\partial a_i^{[1]}}{\partial z_i^{[1]}} \frac{\partial z_i^{[1]}}{\partial W_{ij}^{[1]}} \\ &= (a^{[2]} - y) w_i^{[2]} \left(1 - a_i^{[1]^2}\right) \frac{\partial z_i^{[1]}}{\partial W_{ij}^{[1]}} \\ &= (a^{[2]} - y) w_i^{[2]} \left(1 - a_i^{[1]^2}\right) x_j \end{aligned}$$

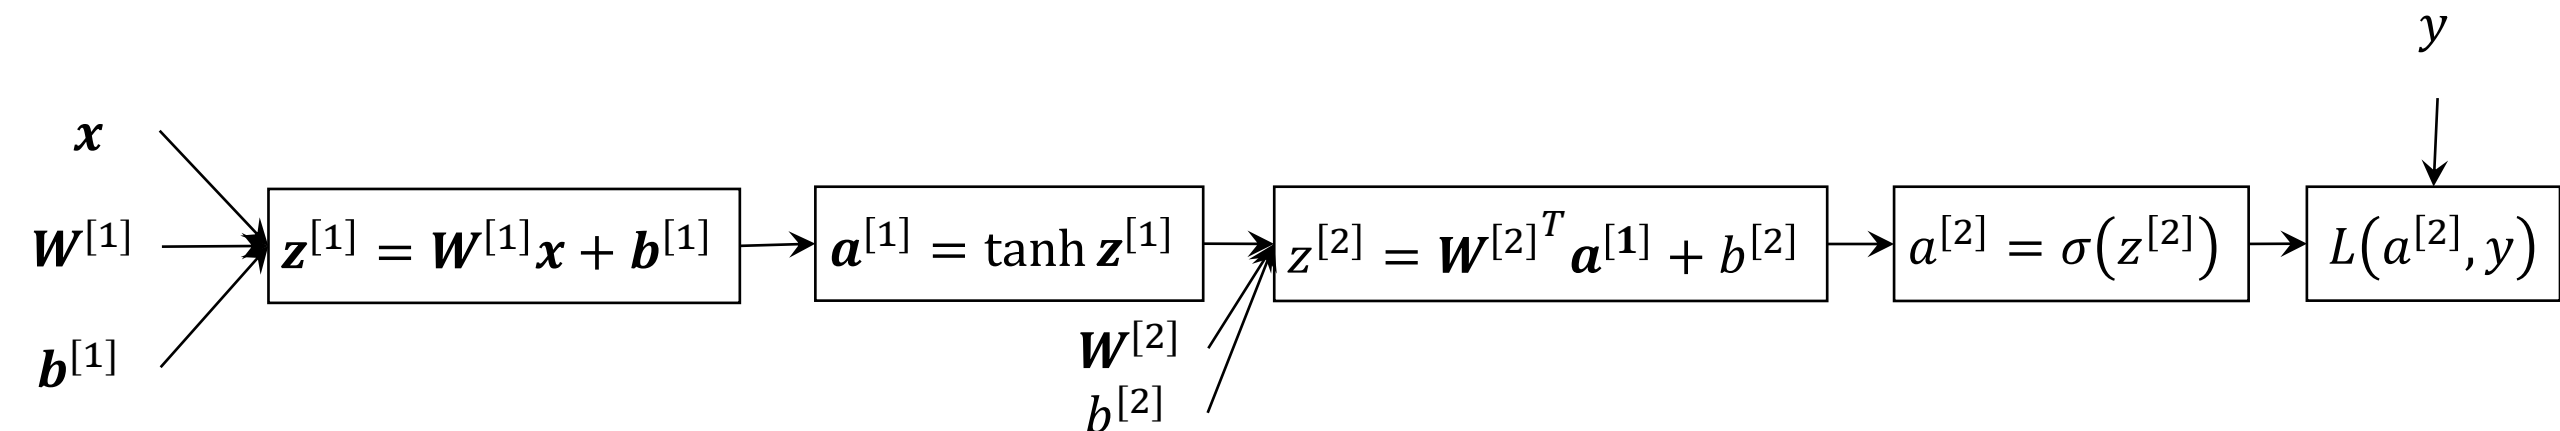
# Partial derivatives

- $\frac{\partial L(a^{[2]}, y)}{\partial b^{[2]}} = a^{[2]} - y$
- $\frac{\partial L(a^{[2]}, y)}{\partial w_i^{[2]}} = (a^{[2]} - y) a_i^{[1]}$
- $\frac{\partial L(a^{[2]}, y)}{\partial b_i^{[1]}} = (a^{[2]} - y) w_i^{[2]} \left(1 - a_i^{[1]^2}\right)$
- $\frac{\partial L(a^{[2]}, y)}{\partial w_{ij}^{[1]}} = (a^{[2]} - y) w_i^{[2]} \left(1 - a_i^{[1]^2}\right) x_j$

# Gradient Descent

- Let  $\mathbf{x} \in \mathbb{R}^n$
- We have  $h$  hidden units
- **Randomly** initialize the parameters
  - $\mathbf{W}^{[1]} \in \mathbb{R}^{h \times n}$
  - $\mathbf{b}^{[1]} \in \mathbb{R}^h$
  - $\mathbf{w}^{[2]} \in \mathbb{R}^h$
  - $b^{[2]} \in \mathbb{R}$
- For each epoch
  - $d\mathbf{W}^{[1]} = \mathbf{0} \in \mathbb{R}^{h \times n}$ ,  $d\mathbf{b}^{[1]} = \mathbf{0} \in \mathbb{R}^h$
  - $d\mathbf{w}^{[2]} = \mathbf{0} \in \mathbb{R}^h$ ,  $db^{[2]} = 0 \in \mathbb{R}$
  - For  $(x, y) \in D$ 
    - $dw_i^{[2]} += \frac{\partial L(a^{[2]}, y)}{\partial w_i^{[2]}}$  for  $1 \leq i \leq h$
    - $db^{[2]} += \frac{\partial L(a^{[2]}, y)}{\partial b^{[2]}}$
    - $dW_{ij}^{[1]} += \frac{\partial L(a^{[2]}, y)}{\partial w_{ij}^{[1]}}$  for  $1 \leq i \leq n$ ,  $1 \leq j \leq h$
    - $db_i^{[1]} += \frac{\partial L(a^{[2]}, y)}{\partial b_i^{[1]}}$  for  $1 \leq i \leq h$ ,
  - $\mathbf{W}^{[1]} -= \eta \cdot d\mathbf{W}^{[1]} / |D|$
  - $\mathbf{b}^{[1]} -= \eta \cdot d\mathbf{b}^{[1]} / |D|$
  - $\mathbf{w}^{[2]} -= \eta \cdot d\mathbf{w}^{[2]} / |D|$
  - $b^{[2]} -= \eta \cdot db^{[2]} / |D|$

# Initialize with the same value?



$$\frac{\partial L(a^{[2]}, y)}{\partial b_i^{[1]}} = \left(1 - a_i^{[1]^2}\right) w_i^{[2]} (a^{[2]} - y)$$

$$\frac{\partial L(a^{[2]}, y)}{\partial b_1^{[1]}} = \frac{\partial L(a^{[2]}, y)}{\partial b_2^{[1]}} = \frac{\partial L(a^{[2]}, y)}{\partial b_3^{[1]}} = \dots$$

# Programming Exercise

Shallow neural network for XOR



# Slicing a numpy array

## Slicing

```
In [77]: a = np.array([2,4,6,8,10])  
         b = np.array([1,3])
```

```
In [78]: a[b]
```

```
Out [78]: array([4, 8])
```

---

# NumPy where

- `numpy.where(condition[, x, y])`
  - Parameters
    - Condition: `array_like`, `bool`
    - x,y: `array_like`
  - Returns:
    - Out: `ndarray`
      - An array with elements from `x` where *condition* is `True`, and elements from `y` elsewhere.

## Examples

```
>>> a = np.arange(10)
>>> a
array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
>>> np.where(a < 5, a, 10*a)
array([ 0,  1,  2,  3,  4, 50, 60, 70, 80, 90])
```

# Slicing a numpy array with condition

```
In [84]: a = np.array([2,4,6,8,10,12])  
        b = np.array([1,3])
```

```
In [85]: a[a>5]
```

```
Out [85]: array([ 6,  8, 10, 12])
```

```
In [86]: a[a%3==0]
```

```
Out [86]: array([ 6, 12])
```

# Data preparation

## XOR data

```
In [4]: x_seeds = np.array([(0,0),(1,0),(0,1),(1,1)],dtype=np.float)
        y_seeds = np.array([0,1,1,0])
```

```
In [5]: N = 1000
        idxs = np.random.randint(0,4,N)
```

```
In [6]: X = x_seeds[idxs]
        Y = y_seeds[idxs]
```

```
In [7]: X += np.random.normal(scale = 0.25, size = X.shape)
```

# Data plotting

- import matplotlib.pyplot as plt


## Plot data

```
In [8]: idxs_1 = np.where(Y==1)
        idxs_0 = np.where(Y==0)
```

```
In [9]: X_0 = X[idxs_0]
        Y_0 = Y[idxs_0]
```

```
In [10]: X_1 = X[idxs_1]
        Y_1 = Y[idxs_1]
```

```
In [11]: #plt.clf()
        plt.plot(X_0[:,0],X_0[:,1],"r^")
        plt.plot(X_1[:,0],X_1[:,1],"bx")
        plt.show()
```



# Model

```
class shallow_neural_network():  
    def __init__(self, num_input_features, num_hidden):  
        self.num_input_features = num_input_features  
        self.num_hidden = num_hidden  
  
        self.W1 = np.random.normal(size = (num_hidden, num_input_features))  
        self.b1 = np.random.normal(size = num_hidden)  
        self.W2 = np.random.normal(size = num_hidden)  
        self.b2 = np.random.normal(size = 1)  
  
    def sigmoid(self, z):  
        return 1/(1 + np.exp(-z))  
  
    def predict(self, x):  
        z1 = np.matmul(self.W1, x) + self.b1  
        a1 = np.tanh(z1)  
        z2 = np.matmul(self.W2, a1) + self.b2  
        a2 = self.sigmoid(z2)  
        return a2, (z1, a1, z2, a2)
```

```
model = shallow_neural_network(2, 3)
```

# Train

## Train

```
] def train(X, Y, model, lr = 0.1):
    dW1 = np.zeros_like(model.W1)
    db1 = np.zeros_like(model.b1)
    dW2 = np.zeros_like(model.W2)
    db2 = np.zeros_like(model.b2)
    m = len(X)
    cost = 0.0
    for x, y in zip(X, Y):
        a2, (z1, a1, z2, _) = model.predict(x)
        if y == 1:
            cost -= np.log(a2)
        else:
            cost -= np.log(1-a2)

        diff = a2-y
        # layer 2
        # db2
        db2 += diff

        # dw2 - todo: remove for-loops
        for i in range(model.num_hidden):
            dW2[i] += a1[i]*diff
        #layer 1
        # db1 - todo: remove for-loops
        for i in range(model.num_hidden):
            db1[i] += (1-a1[i]**2)*model.W2[i]*diff
        # db2 - todo: remove for-loops
        for i in range(model.num_hidden):
            for j in range(model.num_input_features):
                dW1[i, j] += x[j]*(1-a1[i]**2)*model.W2[i]*diff

    cost /= m
    model.W1 -= lr * dW1/m
    model.b1 -= lr * db1/m
    model.W2 -= lr * dW2/m
    model.b2 -= lr * db2/m

    return cost
```

```
for epoch in range(100):
    cost = train(X, Y, model, 1.0)
    if epoch % 10 == 0:
        print(epoch, cost)
```

```
0 [1.16857084]
10 [0.68480409]
20 [0.65975895]
30 [0.60186859]
40 [0.52126503]
50 [0.4437902]
60 [0.3854394]
70 [0.34575853]
80 [0.31936125]
90 [0.30160166]
```

# Test

## Test

```
In [73]: model.predict((1,1))[0].item()
```

```
Out [73]: 0.07201455939561928
```

```
In [74]: model.predict((1,0))[0].item()
```

```
Out [74]: 0.8872553390238738
```

```
In [75]: model.predict((0,1))[0].item()
```

```
Out [75]: 0.858854303287919
```

```
In [76]: model.predict((0,0))[0].item()
```

```
Out [76]: 0.07550823533589167
```



# 숙제

- Train code에서 for-loops 없이 동작하도록 수정
- Hint
  - Layer 2 (Logistic regression참고)
  - Layer 1

`numpy.outer(a, b, out=None)`

Compute the outer product of two vectors.

Given two vectors,  $a = [a_0, a_1, \dots, a_M]$  and  $b = [b_0, b_1, \dots, b_N]$ , the outer product [1] is:

```
[[a0*b0  a0*b1  ...  a0*bN ]
 [a1*b0      .
 [ ...      .
 [aM*b0      aM*bN ]]
```

## Train

```
] : def train(X, Y, model, lr = 0.1):
    dW1 = np.zeros_like(model.W1)
    db1 = np.zeros_like(model.b1)
    dW2 = np.zeros_like(model.W2)
    db2 = np.zeros_like(model.b2)
    m = len(X)
    cost = 0.0
    for x, y in zip(X, Y):
        a2, (z1, a1, z2, _) = model.predict(x)
        if y == 1:
            cost -= np.log(a2)
        else:
            cost -= np.log(1-a2)

        diff = a2 - y
        # layer 2
        # db2
        db2 += diff

        # dw2 - todo: remove for-loops
        for i in range(model.num_hidden):
            dW2[i] += a1[i] * diff
        # layer 1
        # db1 - todo: remove for-loops
        for i in range(model.num_hidden):
            db1[i] += (1-a1[i]**2) * model.W2[i] * diff
        # db2 - todo: remove for-loops
        for i in range(model.num_hidden):
            for j in range(model.num_input_features):
                dW1[i, j] += x[j] * (1-a1[i]**2) * model.W2[i] * diff

    cost /= m
    model.W1 -= lr * dW1 / m
    model.b1 -= lr * db1 / m
    model.W2 -= lr * dW2 / m
    model.b2 -= lr * db2 / m

    return cost
```