

Reinforcement Learning 2



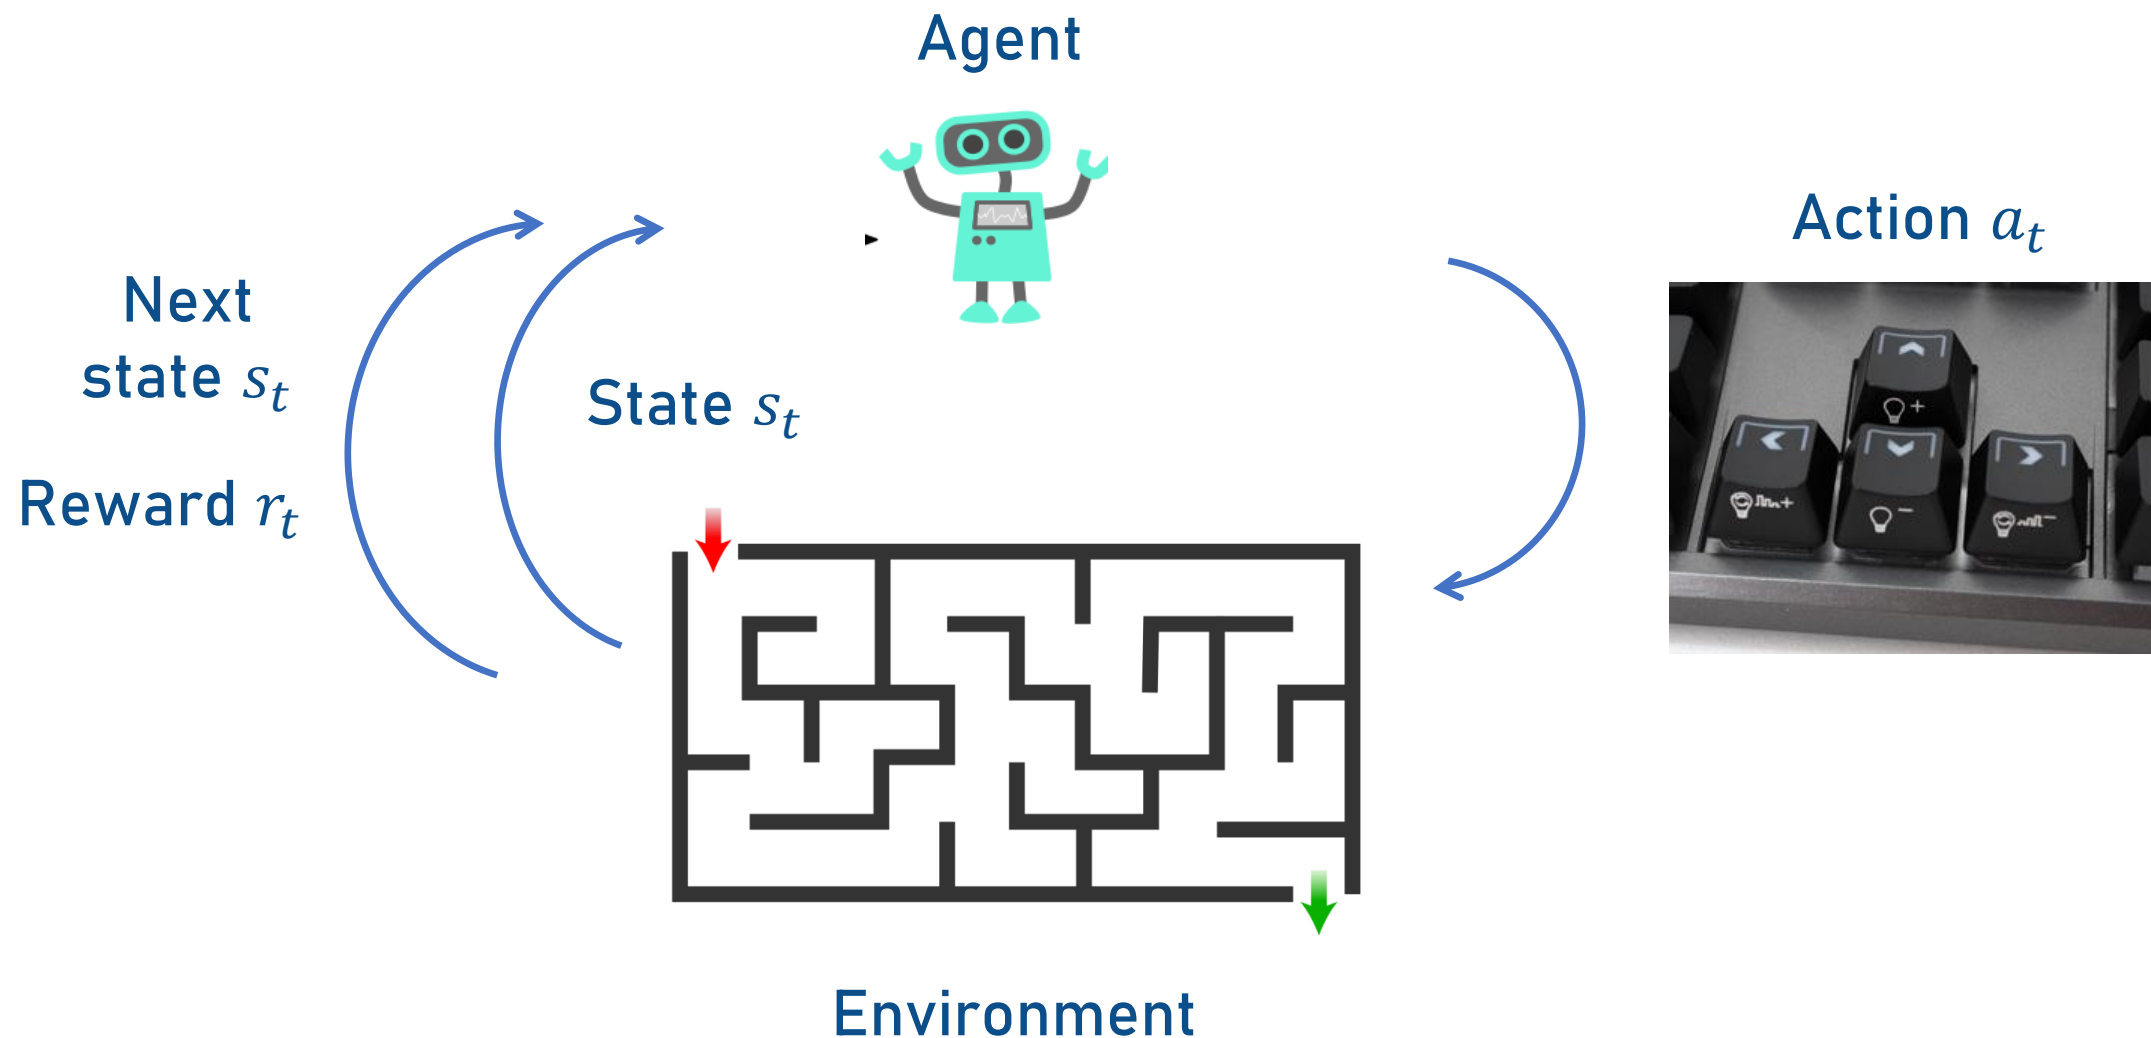
한양대학교 ERICA
소프트웨어융합대학
COLLEGE OF COMPUTING

인공지능학과
Department of
Artificial Intelligence

정 우 환 (whjung@hanyang.ac.kr)

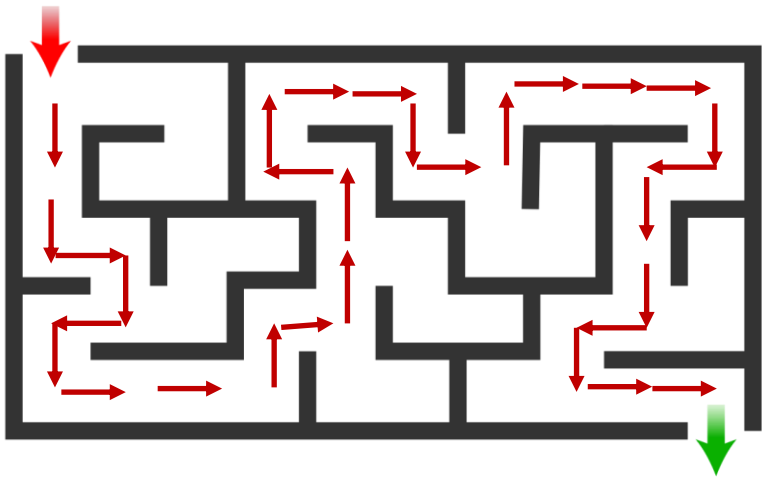
Fall 2021

Reinforcement Learning

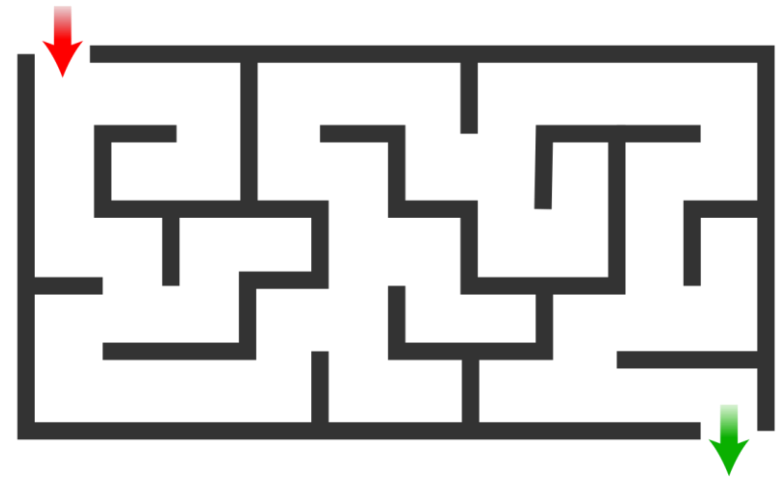


Sparse signal

Labels in supervised learning



Rewards in reinforcement learning



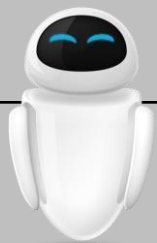
$-T$

Q-learning

Value functions

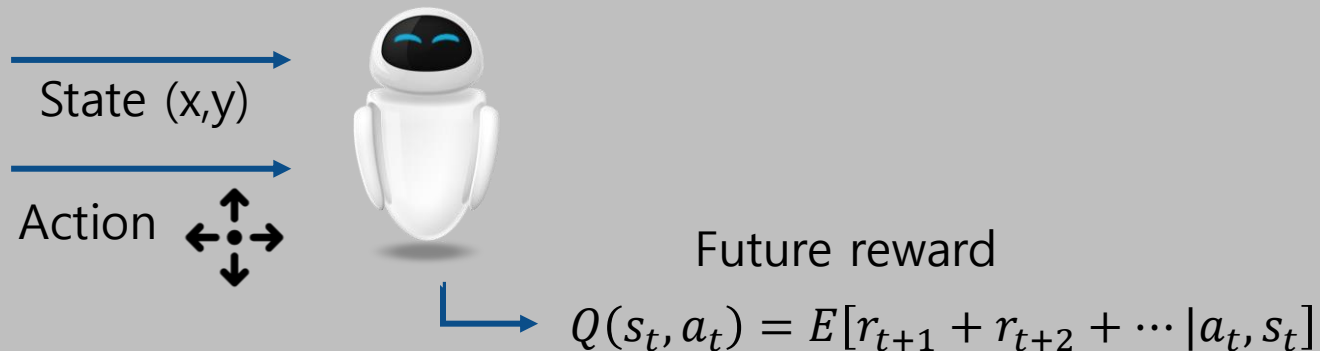
- Total reward $R_t = r_{t+1} + r_{t+2} + \dots$
- Value functions measure the expected total reward after t
 - Value of a **state** $V(s)$
 - $V(s) = E[r_{t+1} + r_{t+2} + \dots | s_t = s]$
 - Value of **of taking an action in a state** $Q(s, a)$
 - $Q(s, a) = E[r_{t+1} + r_{t+2} + \dots | a_t = a, s_t = s]$
- Policy π : mapping from **state** to **action**
 - Optimal policy $\pi^*(s) = \arg \max_a Q(s, a)$

Q-Learning



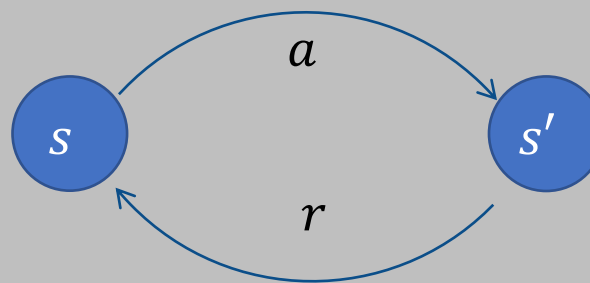
$ \begin{array}{ll} Q((1,1), LEFT) = 0.0 & Q((3,4), LEFT) = 0.0 \\ Q((1,1), RIGHT) = 0.5 & Q((3,4), RIGHT) = 0.0 \\ Q((1,1), UP) = 0.0 & Q((3,4), UP) = 0.0 \\ Q((1,1), DOWN) = 0.3 & Q((3,4), DOWN) = 1.0 \end{array} $			
$ \pi * ((1,1)) \rightarrow RIGHT \quad \pi * ((3,4)) \rightarrow DOWN $			

Q-Function (State-action value) **Q(state,action)**



Optimal policy $\pi^*(s) = \arg \max_a Q(s, a)$

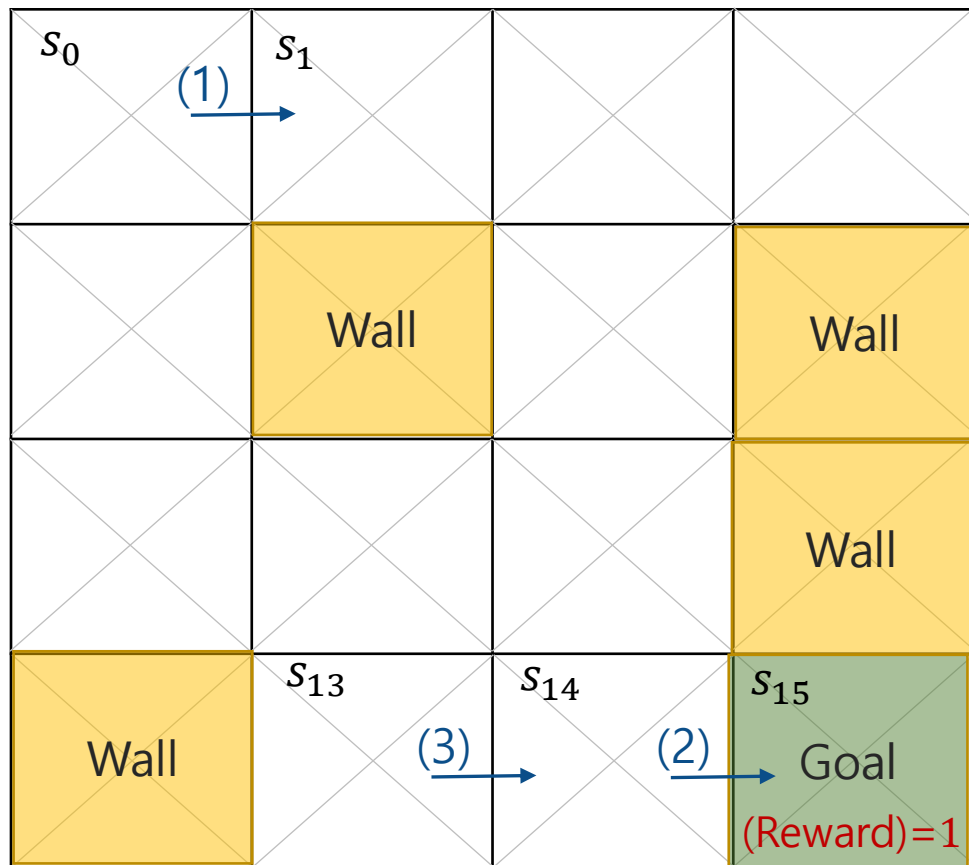
Recurrence equation



$$Q(s, a) = r + \max_{a'} Q(s', a')$$

Q-Learning

16 states and 4 actions (U, D, L, R)



$$Q(s, a) = r + \max_{a'} Q(s', a')$$

- Initial status
 - $Q(s, a) = 0$ for all s, a
 - Reward are all zero except in s_{15}

Case (1) →

$$Q(s_0, R) = r + \max_{a'} Q(s_1, a') = 0 + \max_{a'} \{0, 0, 0, 0\} = 0$$

Case (2) →

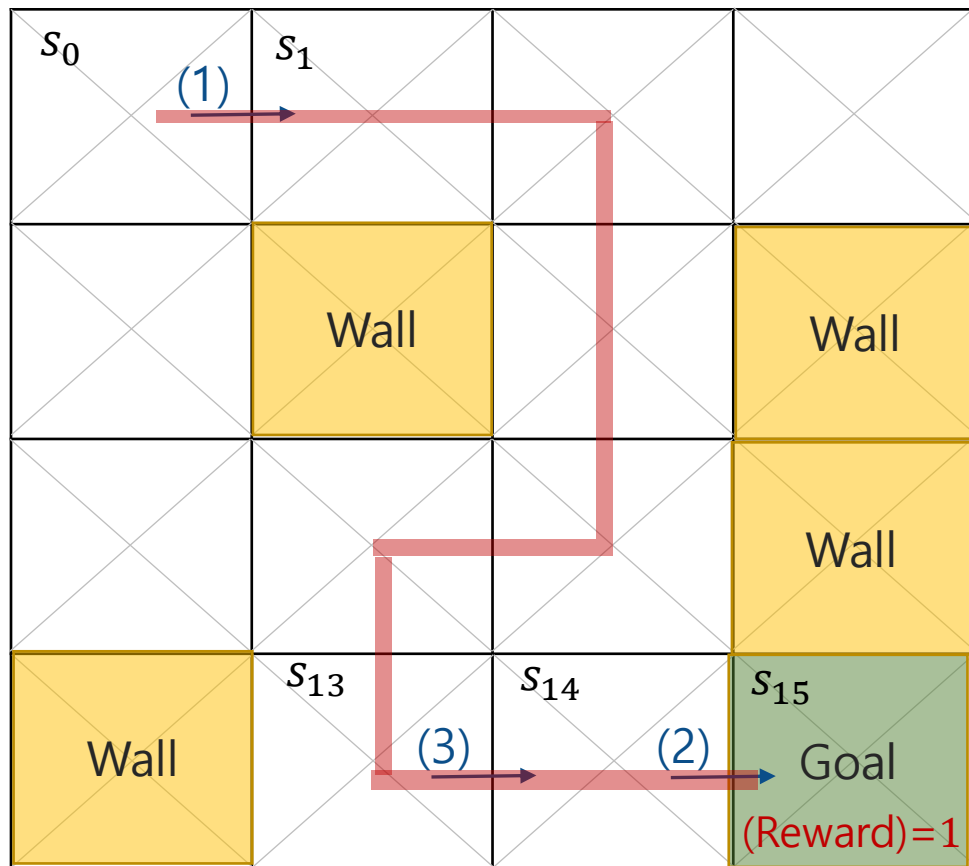
$$Q(s_{14}, R) = 1 + \max_{a'} Q(s_{15}, a') = 0 + \max_{a'} \{0, 0, 0, 0\} = 1$$

Case (3) →

$$Q(s_{13}, R) = r + \max_{a'} Q(s_{14}, a') = 0 + \max_{a'} \{0, 0, 1, 0\} = 1$$

Q-Learning

16 states and 4 actions (U, D, L, R)



$$Q(s, a) = r + \max_{a'} Q(s', a')$$

- Initial status
 - $Q(s, a) = 0$ for all s, a
 - Reward are all zero except in s_{15}

Case (1) →

$$Q(s_0, R) = r + \max_{a'} Q(s_1, a') = 0 + \max_{a'} \{0, 0, 0, 0\} = 0$$

Case (2) →

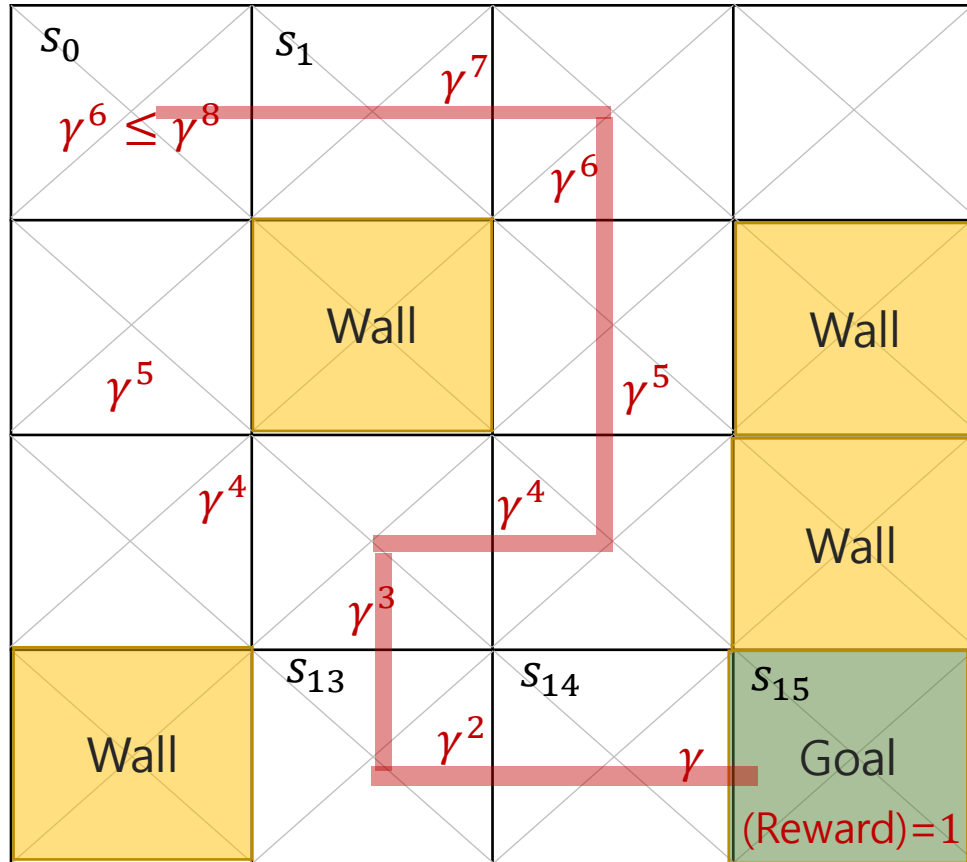
$$Q(s_{14}, R) = 1 + \max_{a'} Q(s_{15}, a') = 0 + \max_{a'} \{0, 0, 0, 0\} = 1$$

Case (3) →

$$Q(s_{13}, R) = r + \max_{a'} Q(s_{14}, a') = 0 + \max_{a'} \{0, 0, 1, 0\} = 1$$

Q-Learning: Discounted reward

16 states and 4 actions (U, D, L, R)



$$Q(s, a) = r + \gamma \max_{a'} Q(s', a')$$

Update with discount factor γ
 $Q(s, a) \leftarrow r + \gamma \max_{a'} Q(s', a')$

Q-Learning: Temporal Difference

$$Q^{new}(s_t, a_t) \leftarrow \underbrace{Q(s_t, a_t)}_{\text{old value}} + \underbrace{\alpha}_{\text{learning rate}} \cdot \underbrace{\left(\underbrace{r_t}_{\text{reward}} + \underbrace{\gamma}_{\text{discount factor}} \cdot \underbrace{\max_a Q(s_{t+1}, a)}_{\text{estimate of optimal future value}} - \underbrace{Q(s_t, a_t)}_{\text{old value}} \right)}_{\text{temporal difference}}$$

new value (temporal difference target)

Range of α : $0 < \alpha < 1$

If learning rate $\alpha = 1$, $Q(s, a) \leftarrow r + \gamma \max_{a'} Q(s', a')$

If learning rate $\alpha = 0$, $Q(s, a) \leftarrow Q(s, a)$

Q-learning

- For each s, a , initialize table entry $Q(s, a) \leftarrow 0$
- Do until Q converges
 - Initialize s
 - Do until s is terminal
 - Select an action a using policy π derived from Q
 - Take action a
 - Receive immediate reward r
 - Observe the new state s'
 - $Q(s, a) \leftarrow Q(s, a) + \alpha \left[r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right]$
 - $s \leftarrow s'$

Iterative update
Discounted reward
Temporal difference

Policy π

- Testing phase
 - Optimal policy $\pi^*(s) = \arg \max_a Q(s, a)$
- Training phase (**ϵ -greedy**)
 - **Exploration-Exploitation**
 - Exploration allows an agent to improve its current knowledge
 - Exploitation chooses the greedy action to get the most reward by exploiting the agent's current action-value estimates

Action at time(t) $\left\{ \begin{array}{ll} \max Q_t(a) & \text{with probability } 1-\epsilon \\ \text{any action (a)} & \text{with probability } \epsilon \end{array} \right.$

Q-learning

- For each s, a , initialize table entry $Q(s, a) \leftarrow 0$
- Do until Q converges
 - Initialize s
 - Do until s is terminal
 - Draw a random value $v \sim \text{Uniform}(0,1)$
 - If $v < \epsilon$
 - Randomly select a
 - Else:
 - $a = \underset{a'}{\operatorname{argmax}} Q(s, a')$
 - Take action a
 - Receive immediate reward r
 - Observe the new state s'
 - $Q(s, a) \leftarrow Q(s, a) + \alpha \left[r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right]$
 - $s \leftarrow s'$

ϵ -greedy

Implementing a Q-learning algorithm in Python

Taxi-v3

<https://www.learndatasci.com/tutorials/reinforcement-q-learning-scratch-python-openai-gym/>



Gym

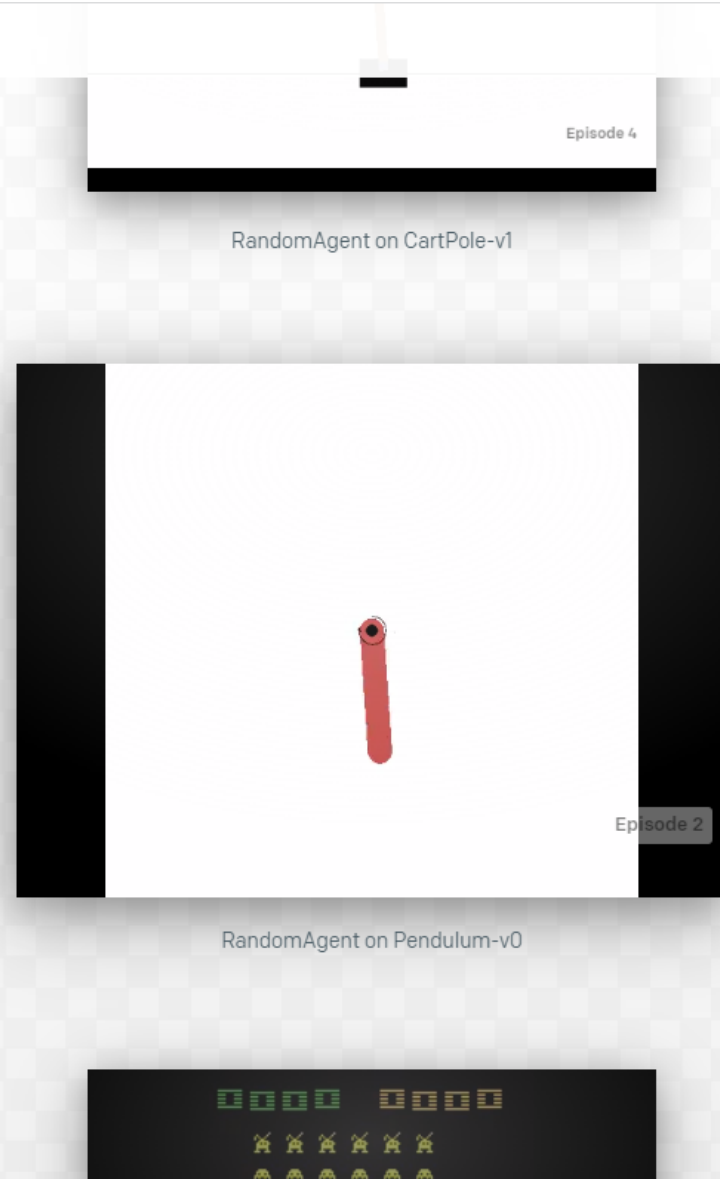
Gym is a toolkit for developing and comparing reinforcement learning algorithms. It supports teaching agents everything from walking to playing games like Pong or Pinball.

[View documentation >](#)

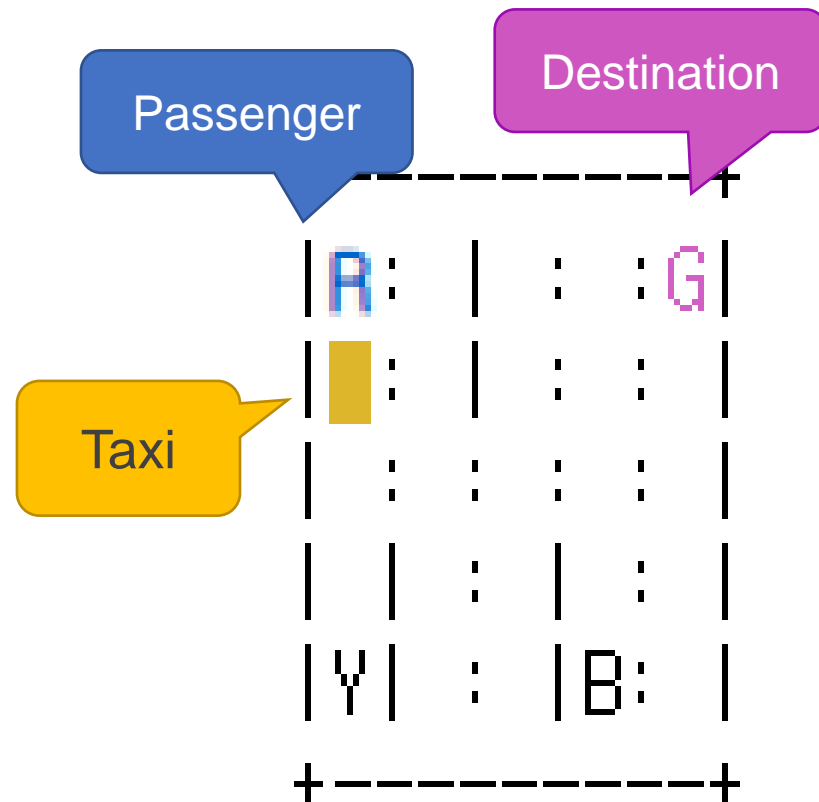
[View on GitHub >](#)



Open source interface to reinforcement learning tasks.
The [gym](#) library provides an easy-to-use suite of reinforcement



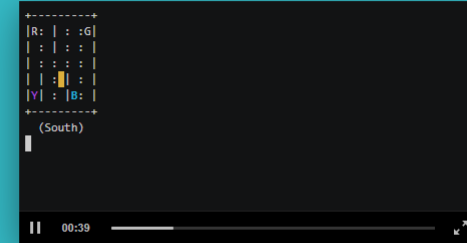
Taxi-v3



Taxi-v3

This task was introduced in [Dietterich2000] to illustrate some issues in hierarchical reinforcement learning. There are 4 locations (labeled by different letters) and your job is to pick up the passenger at one location and drop him off in another. You receive +20 points for a successful dropoff, and lose 1 point for every timestep it takes. There is also a 10 point penalty for illegal pick-up and drop-off actions.

[Dietterich2000] T Erez, Y Tassa, E Todorov, "Hierarchical Reinforcement Learning with the MAXQ Value Function Decomposition", 2011.

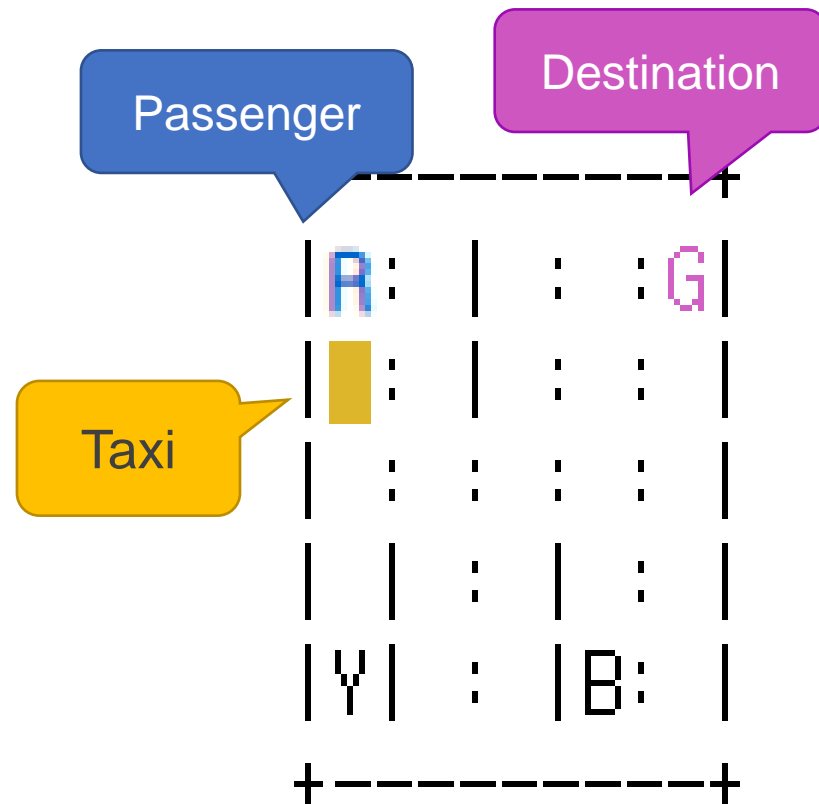


<https://gym.openai.com/envs/Taxi-v3/>

https://github.com/openai/gym/blob/master/gym/envs/toy_text/taxi.py

- 4 designated locations: R(ed), G(reen), Y(ellow), and B(lue)
- When the episode starts
 - The taxi starts off at a random location
 - The passenger is at a random location
- Goal
 - Taxi picks up the passenger and drives to the destination with the minimum operations

Taxi-v3



- 6 actions
 - 0: move south
 - 1: move north
 - 2: move east
 - 3: move west
 - 4: pickup passenger
 - 5: drop off passenger
- States
 - (Taxi row, taxi col, passenger_location, destination)
 - $(5 \times 5 \times (4 + 1) \times 4) = 500$ possible states
 - Passenger location: RGBY + taxi
- Rewards
 - Per step reward: -1
 - Delivering the passenger: +20
 - Illegal pickup and drop-off: -10

Taxi-v3

- Initialize environment

```
In [1]: import gym
```

```
In [2]: env = gym.make("Taxi-v3").env
```

- Render environment

```
In [3]: env.render()
```

```
+-----+
|R: | : :G|
| : | : :|
| : : : :|
| | : | :|
|Y| : |B: |
+-----+
```

```
In [4]: env.reset()
env.render()
```

```
+-----+
|R: | : :G|
| : | : :|
| : : : :|
| : | : :|
|Y| : |B: |
+-----+
```

```
In [5]: env.reset()
env.render()
```

```
+-----+
|R: | : :G|
| : | : :|
| : : : :|
| | : |Y|
|Y| : |B: |
+-----+
```

Action space and state space

```
: print("Action Space {}".format(env.action_space))  
  print("State Space {}".format(env.observation_space))
```

Action Space Discrete(6)
State Space Discrete(500)

States

(Taxi row, taxi col, passenger_location, destination)

```
state = env.encode(3, 1, 2, 0)  
print("State:", state)
```

```
env.s = state  
env.render()
```

State: 328

```
+-----+  
|R: | : :G| |
| : | : :|  
| : | : :|  
| : | : :|  
| |Y: | : :|  
|Y| : |B: |  
+-----+
```

```
state = env.encode(3, 1, 3, 2)  
print("State:", state)
```

```
env.s = state  
env.render()
```

State: 334

```
+-----+  
|R: | : :G| |
| : | : :|  
| : | : :|  
| : | : :|  
| |Y: | : :|  
|Y| : |B: |  
+-----+
```

Reward table P

6 actions

- 0: move south
- 1: move north
- 2: move east
- 3: move west
- 4: pickup passenger
- 5: drop off passenger

```
state = env.encode(3, 1, 2, 0)
print("State:", state)
```

```
env.s = state
env.render()
```

State: 328

```
+-----+
|R: | : :G| |
| : | : :|
| : | : :|
| : | : :|
| | | : :|
|Y| : |B:|
+-----+
```

```
In [22]: env.P[328]
```

```
Out [22]: {0: [(1.0, 428, -1, False)],
           1: [(1.0, 228, -1, False)],
           2: [(1.0, 348, -1, False)],
           3: [(1.0, 328, -1, False)],
           4: [(1.0, 328, -10, False)],
           5: [(1.0, 328, -10, False)]}
```

Brute force algorithm

- Brute-force search is a very general problem-solving technique
 - Enumerates all possible candidates
 - Checks whether each candidate satisfies the problem's statement

```
env.s = 328 # set environment to illustration's state

epochs = 0
penalties, reward = 0, 0

frames = [] # for animation

done = False

while not done:
    action = env.action_space.sample()
    state, reward, done, info = env.step(action)

    if reward == -10:
        penalties += 1

    # Put each rendered frame into dict for animation
    frames.append({
        'frame': env.render(mode='ansi'),
        'state': state,
        'action': action,
        'reward': reward
    })

    epochs += 1

print("Timesteps taken: {}".format(epochs))
print("Penalties incurred: {}".format(penalties))
```

Timesteps taken: 620
Penalties incurred: 208

Brute force algorithm

```
: from IPython.display import clear_output
  from time import sleep

  def print_frames(frames, s = .01):
      for i, frame in enumerate(frames):
          clear_output(wait=True)
          print(frame['frame'])
          print(f"Timestep: {i + 1}")
          print(f"State: {frame['state']}")
          print(f"Action: {frame['action']}")
          print(f"Reward: {frame['reward']}")
          sleep(s)

: print_frames(frames)
```

Q-learning

- Q-table size: (*#states*, *#actions*)

```
import numpy as np
q_table = np.zeros([env.observation_space.n, env.action_space.n])
```

Q-learning: Training

- For each s, a , initialize table entry $Q(s, a) \leftarrow 0$
- Do until Q converges
 - Initialize s
 - Do until s is terminal
 - Draw a random value $v \sim \text{Uniform}(0,1)$
 - If $v < \epsilon$
 - Randomly select a
 - Else:
 - $a = \underset{a'}{\operatorname{argmax}} Q(s, a')$
 - Take action a
 - Receive immediate reward r
 - Observe the new state s'
 - $Q(s, a) \leftarrow Q(s, a) + \alpha \left[r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right]$
 - $s \leftarrow s'$

```
from IPython.display import clear_output

# Hyperparameters
alpha = 0.1
gamma = 0.6
epsilon = 0.1

# For plotting metrics
all_epochs = []
all_penalties = []

for i in range(1, 100001):
    state = env.reset()

    epochs, penalties, reward, = 0, 0, 0
    done = False

    while not done:
        if random.uniform(0, 1) < epsilon:
            action = env.action_space.sample() # Explore action space
        else:
            action = np.argmax(q_table[state]) # Exploit learned values

        next_state, reward, done, info = env.step(action)

        old_value = q_table[state, action]
        next_max = np.max(q_table[next_state])

        new_value = (1 - alpha) * old_value + alpha * (reward + gamma * next_max)
        q_table[state, action] = new_value

        if reward == -10:
            penalties += 1

        state = next_state
        epochs += 1

    if i % 100 == 0:
        clear_output(wait=True)
        print(f"Episode: {i}")
```


Q-learning: Test

$$\pi^*(s) = \arg \max_a Q(s, a)$$

```
total_epochs, total_penalties = 0, 0
episodes = 100

for _ in range(episodes):
    state = env.reset()
    epochs, penalties, reward = 0, 0, 0

    done = False

    while not done:
        action = np.argmax(q_table[state])
        state, reward, done, info = env.step(action)

        if reward == -10:
            penalties += 1

        epochs += 1

    total_penalties += penalties
    total_epochs += epochs

print(f"Results after {episodes} episodes:")
print(f"Average timesteps per episode: {total_epochs / episodes}")
print(f"Average penalties per episode: {total_penalties / episodes}")
```

```
Results after 100 episodes:
Average timesteps per episode: 13.23
Average penalties per episode: 0.0
```

Q-learning: Test (print result)

```
frames = []

state = env.reset()
epochs, penalties, reward = 0, 0, 0

done = False

while not done:
    action = np.argmax(q_table[state])
    state, reward, done, info = env.step(action)

    if reward == -10:
        penalties += 1

    epochs += 1
    frames.append({
        'frame': env.render(mode='ansi'),
        'state': state,
        'action': action,
        'reward': reward
    })

total_penalties += penalties
total_epochs += epochs
```

```
print_frames(frames, s = 0.1)
```

```
+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+
      (Dropoff)
```

```
Timestep: 12
State: 85
Action: 5
Reward: 20
```

References

- Fei-Fei Li & Justin Johnson & Serena Yeung, CS231n Lecture 14: Reinforcement Learning, Stanford University
- <https://sumniya.tistory.com/> **숨니의 무작정 따라하기**
- <https://youtu.be/m1FC3dMmY78> Joongheon Kim, Korea University
- <https://www.learndatasci.com/tutorials/reinforcement-q-learning-scratch-python-openai-gym/>