


Logistic Regression

Deep Learning

Woohwan Jung



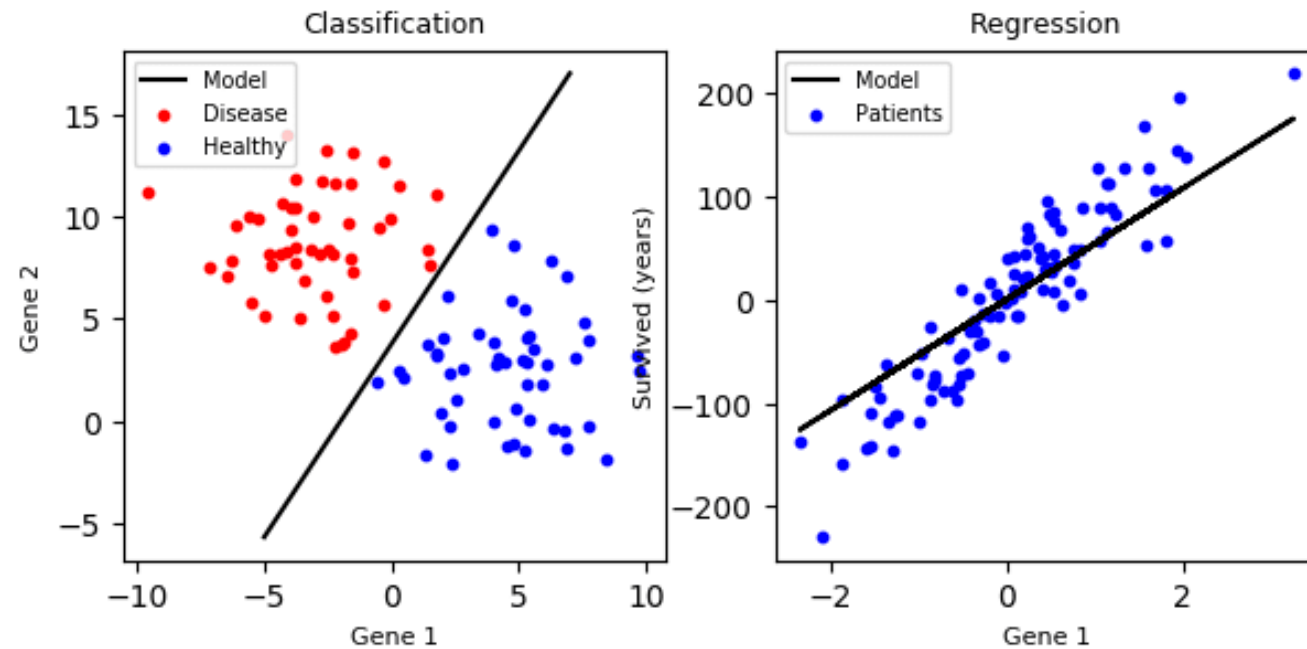
오픈카톡방

- <https://open.kakao.com/o/gCLPx6Ce>
- 수업중
 - 말로하기 어려운 복잡한 질문은 카톡으로 (말로하는걸 권장)
- 수업시간외
 - 학생들간 의견 및 자료교환
 - 학생들간 질문 (오픈카톡방에서 공유된 내용은 Copy X)

Topics


- Binary Classification
- Logistic Regression
 - Model
 - Cost/Loss function
 - Optimization

Classification vs Regression



Logistic regression algorithm is used for **(binary) classification**!

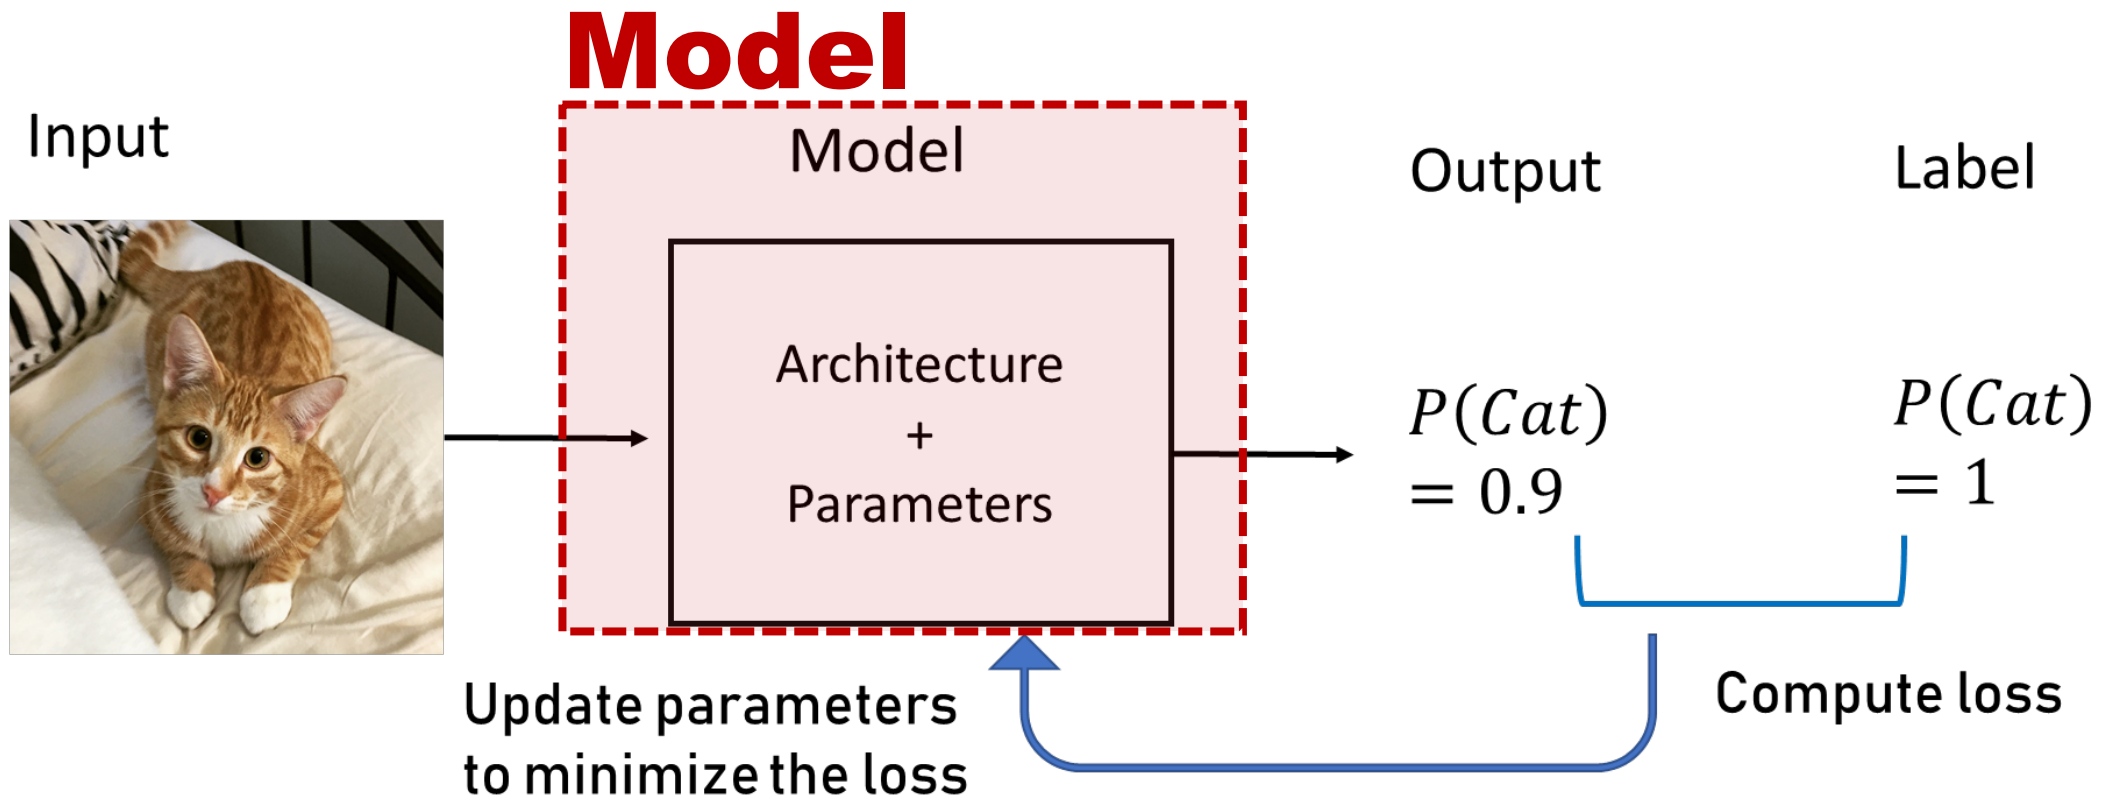
Binary Classification



1(cat) vs 0 (non cat)

		Blue			
	Green	123	94	83	2
Red		123	94	83	4
		123	94	83	2
		34	44	187	92
		34	76	232	124
		67	83	194	202

$$x = \begin{bmatrix} 123 \\ 94 \\ \dots \\ 202 \\ 123 \\ 94 \\ \dots \\ 142 \end{bmatrix}$$



Logistic Regression

Logistic Regression

- A simple model for **binary classification**
- Maybe one of the simplest neural network
- A training example (\mathbf{x}, y)
 - Input: $\mathbf{x} \in \mathbb{R}^n$
 - Output: $y \in \{0,1\}$
- m training examples
 - $\{(\mathbf{x}^{(1)}, y^{(1)}), (\mathbf{x}^{(2)}, y^{(2)}), \dots, (\mathbf{x}^{(m)}, y^{(m)})\}$

Model = Architecture
+ Parameters

Review: Linear Regression

- Given $\mathbf{x} \in \mathbb{R}^n$
- Want $\hat{y} \approx y$

$$\hat{y} = \mathbf{w}^\top \mathbf{x} + b$$

Parameters: $\mathbf{w} \in \mathbb{R}^n$, $b \in \mathbb{R}$

Model = Architecture
+ Parameters

Logistic Regression

- Given $\mathbf{x} \in \mathbb{R}^n$
- Want $\hat{y} = P(y = 1|\mathbf{x})$

$$\hat{y} = \sigma(\mathbf{w}^\top \mathbf{x} + b)$$

Parameters: $\mathbf{w} \in \mathbb{R}^n$, $b \in \mathbb{R}$

$$\text{Sigmoid } \sigma(x) = \frac{1}{1+e^{-x}}$$

$$\sigma(-\infty) = 0$$

$$\sigma(+\infty) = 1$$

$$\Theta = \{\mathbf{w}, b\}$$

Model = Architecture
+ Parameters


Logistic Regression

- Given $\mathbf{x} \in \mathbb{R}^n$
- Want $\hat{y} = P(y = 1|\mathbf{x})$

$$\hat{y} = \sigma(\mathbf{w}^\top \mathbf{x} + b)$$

Parameters: $\mathbf{w} \in \mathbb{R}^n$, $b \in \mathbb{R}$


$$\text{Sigmoid } \sigma(x) = \frac{1}{1+e^{-x}}$$



Input



Model



Output

Label

$$P(\text{Cat}) = 0.9$$

$$P(\text{Cat}) = 1$$

Compute loss

Update parameters
to minimize the loss

Logistic Regression

**Cost function
& Loss**

Logistic Regression: Cost function

Roughly, $\hat{y}^{(i)} \approx y^{(i)}$

- $\hat{y} = \sigma(\mathbf{w}^\top \mathbf{x} + b)$ where $\sigma(z) = \frac{1}{1+e^{-z}}$
- Given $\{(\mathbf{x}^{(1)}, y^{(1)}), \dots (\mathbf{x}^{(m)}, y^{(m)})\}$, want $\hat{y}^{(i)} = P(y^{(i)} = 1)$
- Loss function: Binary Cross Entropy

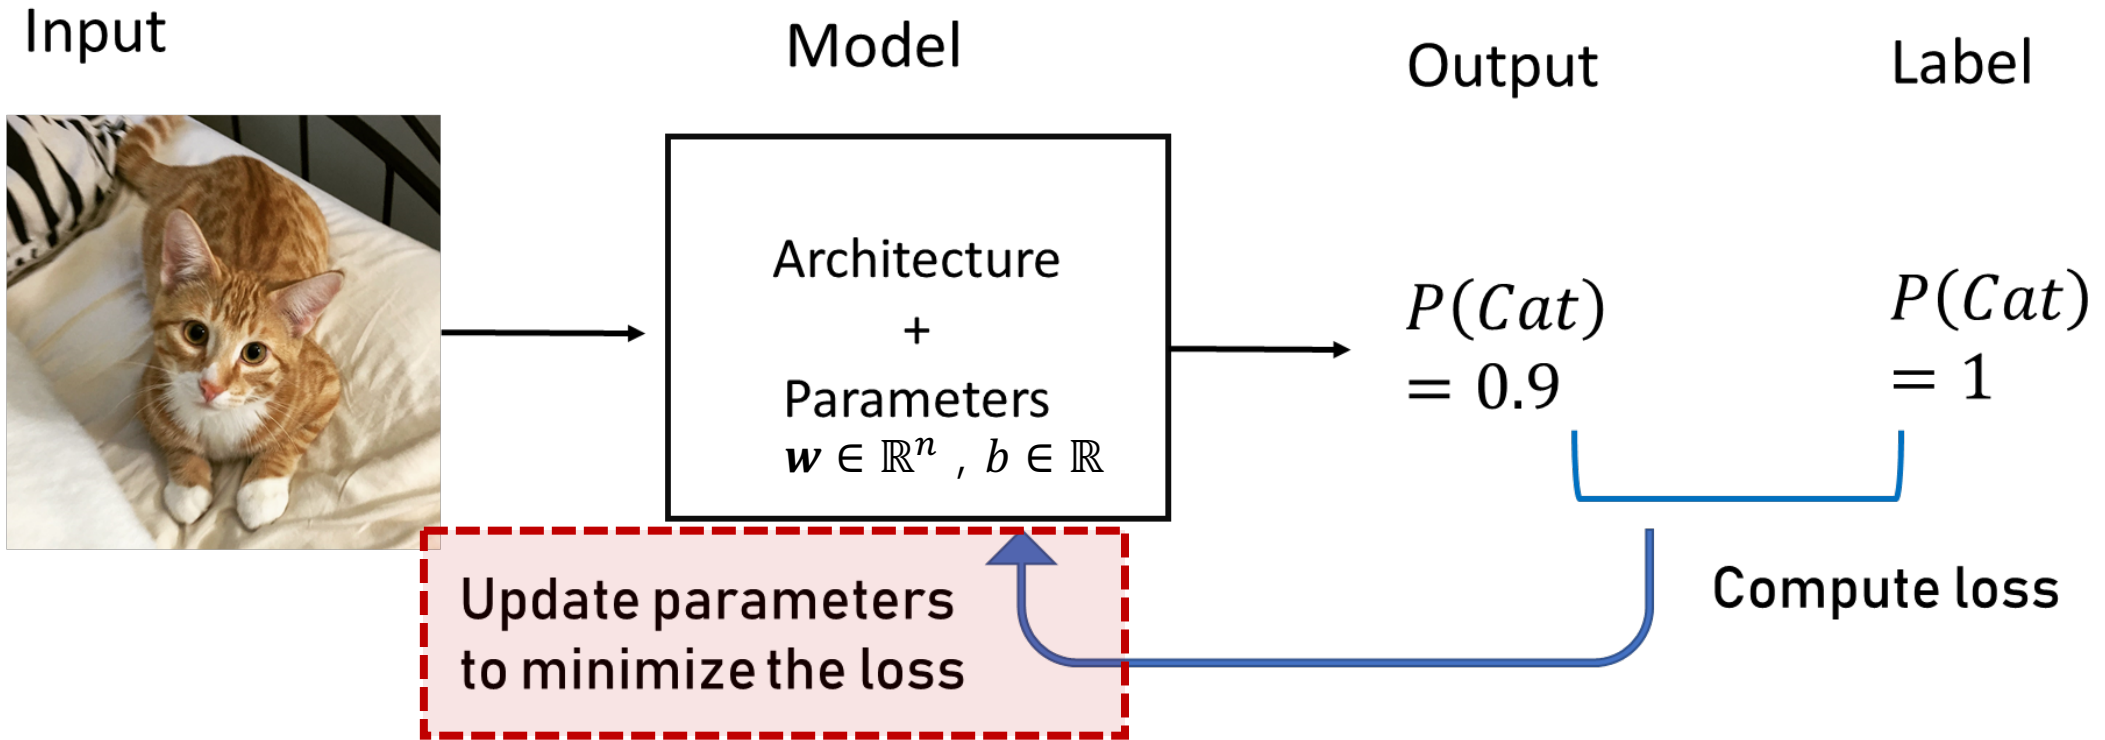
$$L(\hat{y}, y) = -y \log \hat{y} - (1 - y) \log(1 - \hat{y})$$

$$\text{If } y = 1: L(\hat{y}, y) = -\log \hat{y}$$

$$\text{If } y = 0: L(\hat{y}, y) = -\log(1 - \hat{y})$$

- Cost function:

$$J(\mathbf{w}, b) = \frac{1}{m} \sum_{i=1}^m L(\hat{y}^{(i)}, y^{(i)})$$



Logistic Regression

Optimization

Optimization

- Logistic regression model

- $\hat{y} = \sigma(\mathbf{w}^\top \mathbf{x} + b)$ where $\sigma(z) = \frac{1}{1+e^{-z}}$

- Cost function

$$J(\mathbf{w}, b) = \frac{1}{m} \sum_{i=1}^m L(\hat{y}^{(i)}, y^{(i)}) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})]$$

- Our goal

- Find parameters $\mathbf{w} \in \mathbb{R}^n$, $b \in \mathbb{R}$ that minimize $J(\mathbf{w}, b)$


- Gradient Descent!

Gradient Descent

- Algorithm to minimize a cost function $J(\theta)$
- $J(\theta)$: cost function
- θ : model parameters
- η : Learning rate

Repeatedly update


$$\theta = \theta - \eta \cdot \nabla_{\theta} J(\theta)$$



Derivative with a Computation Graph

Computation Graph

■ $J(a, b, c) = 3(a + bc)$

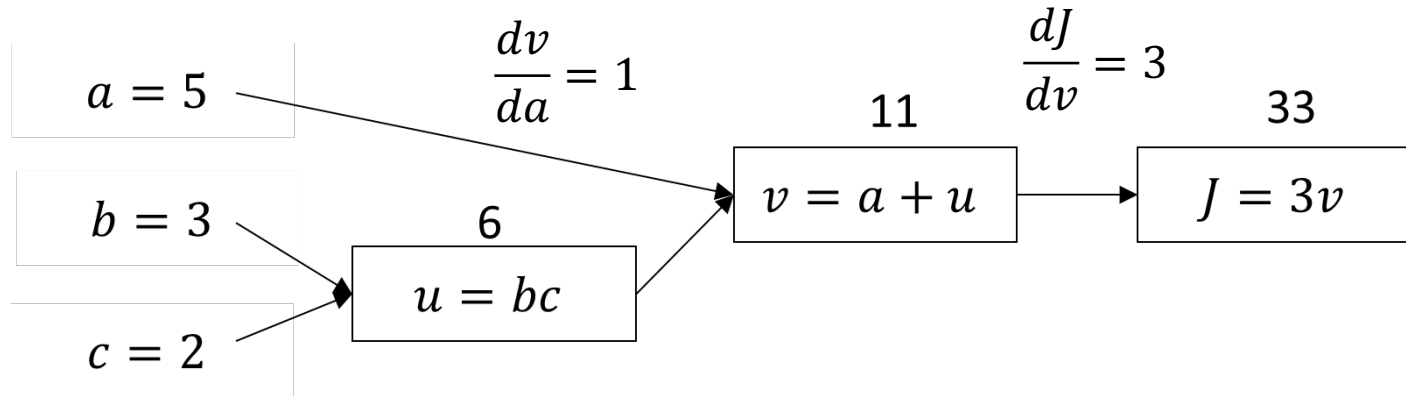


Chain Rule (calculus)

$$(f \circ g)'(c) = f'(g(c)) \cdot g'(c)$$

$$\frac{dy}{dx} = \frac{dy}{du} \cdot \frac{du}{dx}$$

Derivatives with a Computation Graph




$$\frac{dy}{dx} = \frac{dy}{du} \cdot \frac{du}{dx}$$

Chain rule

$$\frac{dJ}{da} = \frac{dJ}{dv} \frac{dv}{da} =$$

Derivatives with a Computation Graph



$$\frac{dy}{dx} = \frac{dy}{du} \cdot \frac{du}{dx}$$

Chain rule

$$\frac{dJ}{da} = \frac{dJ}{dv} \frac{dv}{da} = 3$$

$$\frac{dJ}{db} = \frac{dJ}{du} \frac{du}{db} = 2 \quad \frac{dJ}{du} = 2 \quad \frac{dv}{du} \frac{dJ}{dv} = 6$$

$$\frac{dJ}{dc} = \frac{dJ}{du} \frac{du}{dc} = 3 \quad \frac{dJ}{du} = 3 \quad \frac{dv}{du} \frac{dJ}{dv} = 9$$

Gradient Descent :Logistic Regression

Background: Derivatives

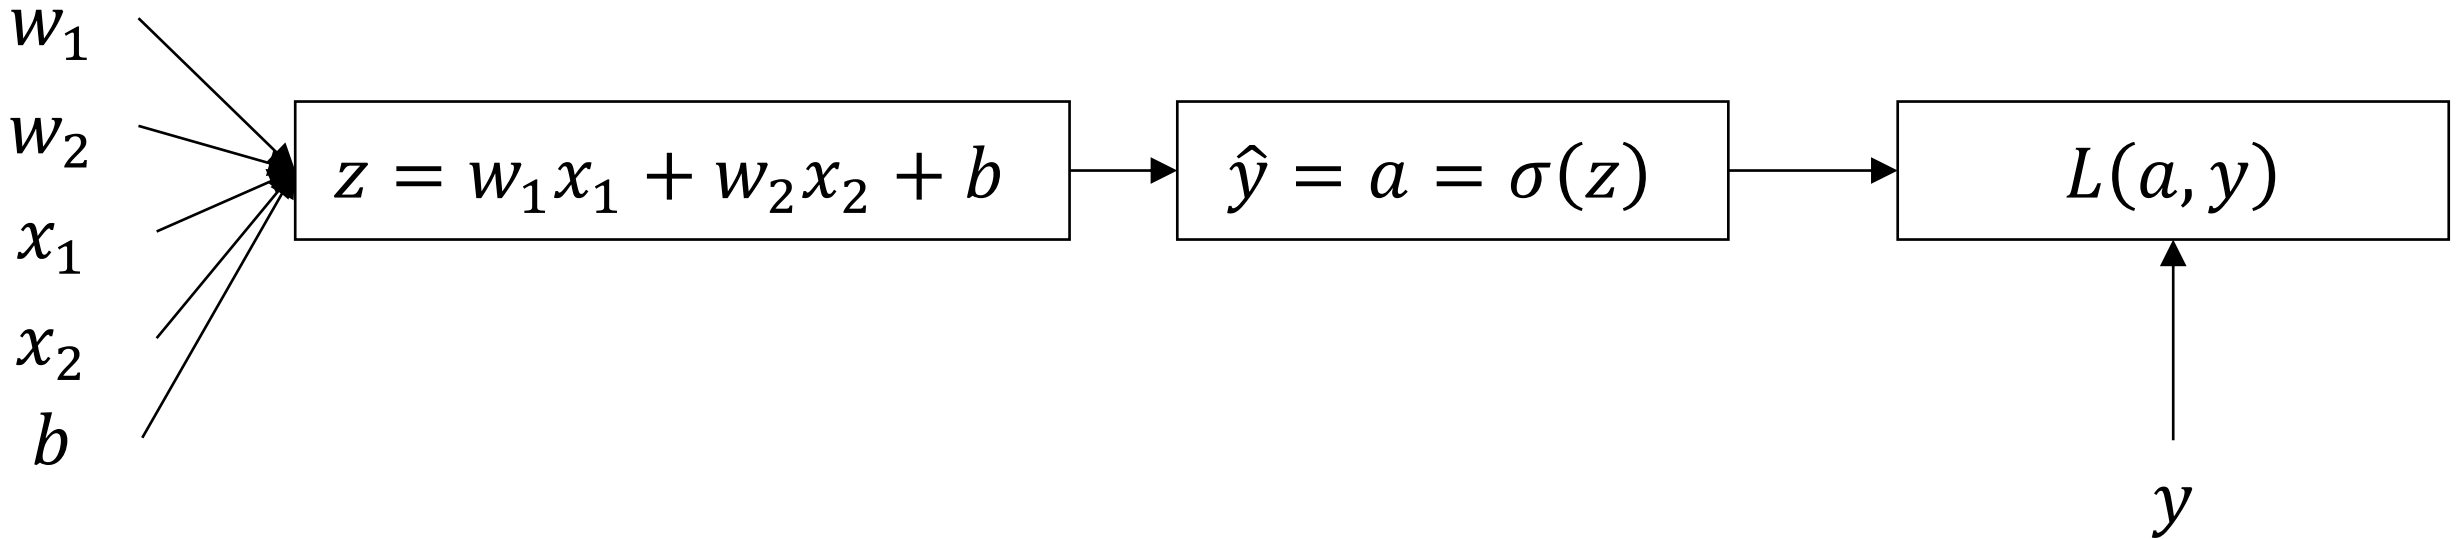
$$1. \frac{d}{dx} \log_e x = \frac{1}{x}$$

$$2. \frac{d}{dx} \sigma(x) = \sigma(x)(1 - \sigma(x))$$

Logistic Regression Recap

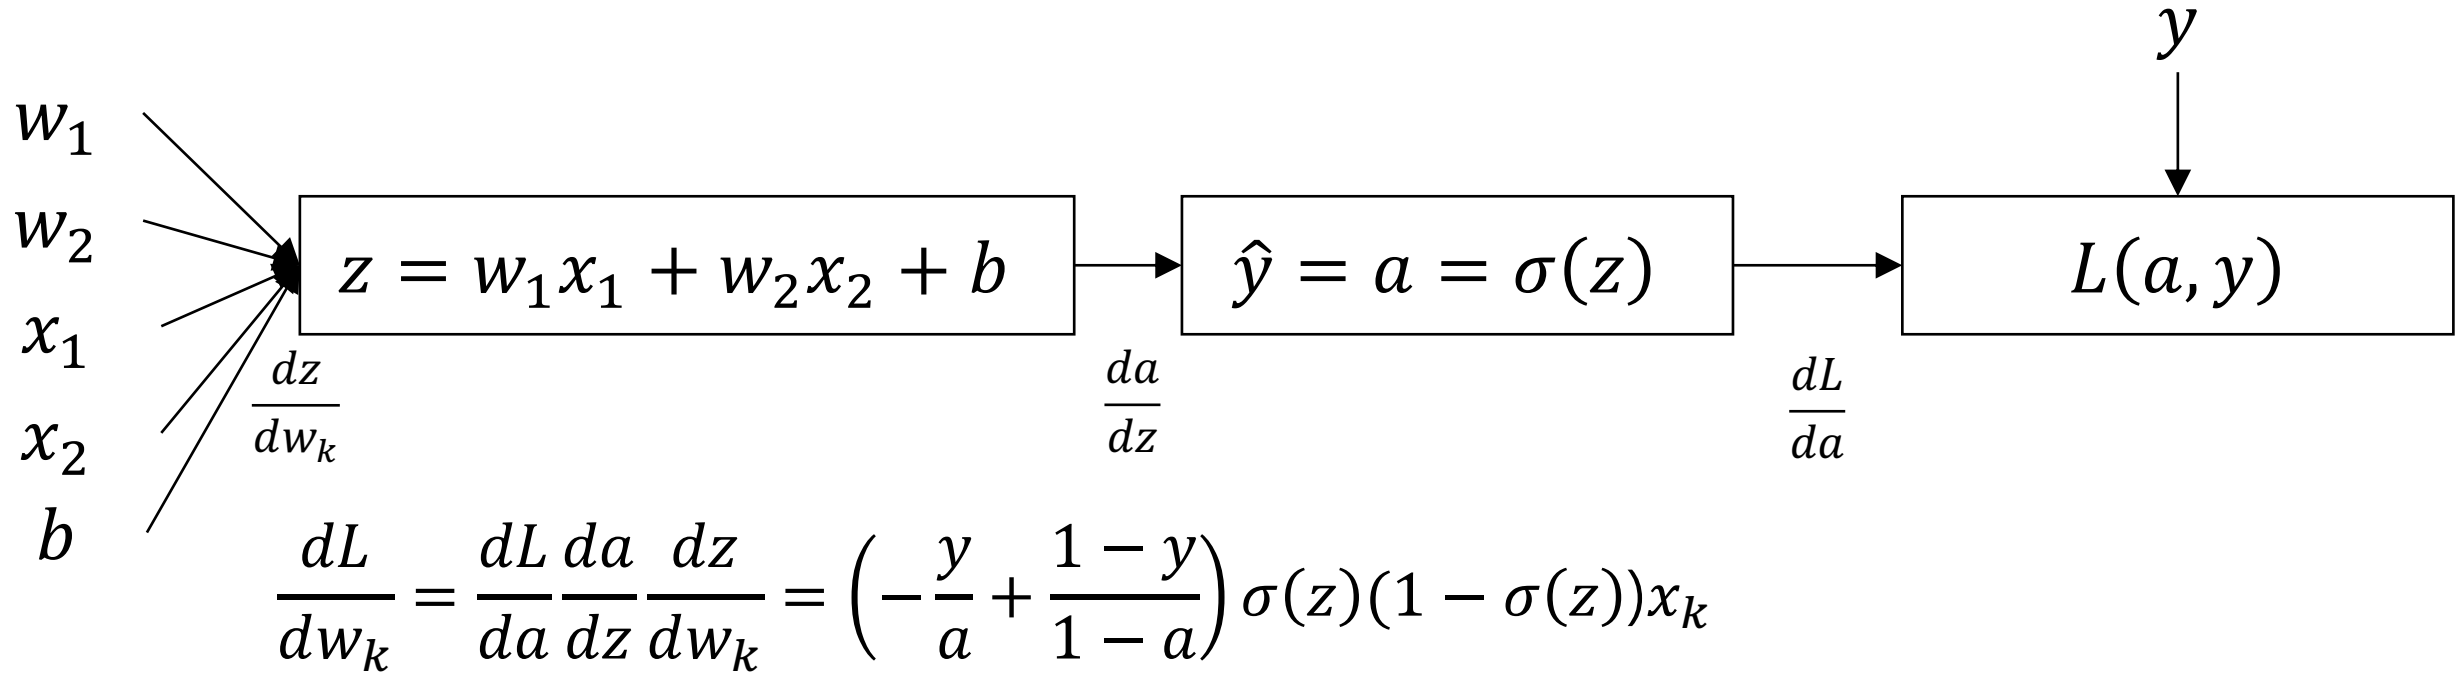
- $\hat{y} = \sigma(\mathbf{w}^\top \mathbf{x} + b)$ where $\sigma(z) = \frac{1}{1+e^{-z}}$
- $L(\hat{y}, y) = -y \log \hat{y} - (1 - y) \log(1 - \hat{y})$

For the simplicity
 $\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$, $\mathbf{w} = \begin{bmatrix} w_1 \\ w_2 \end{bmatrix}$



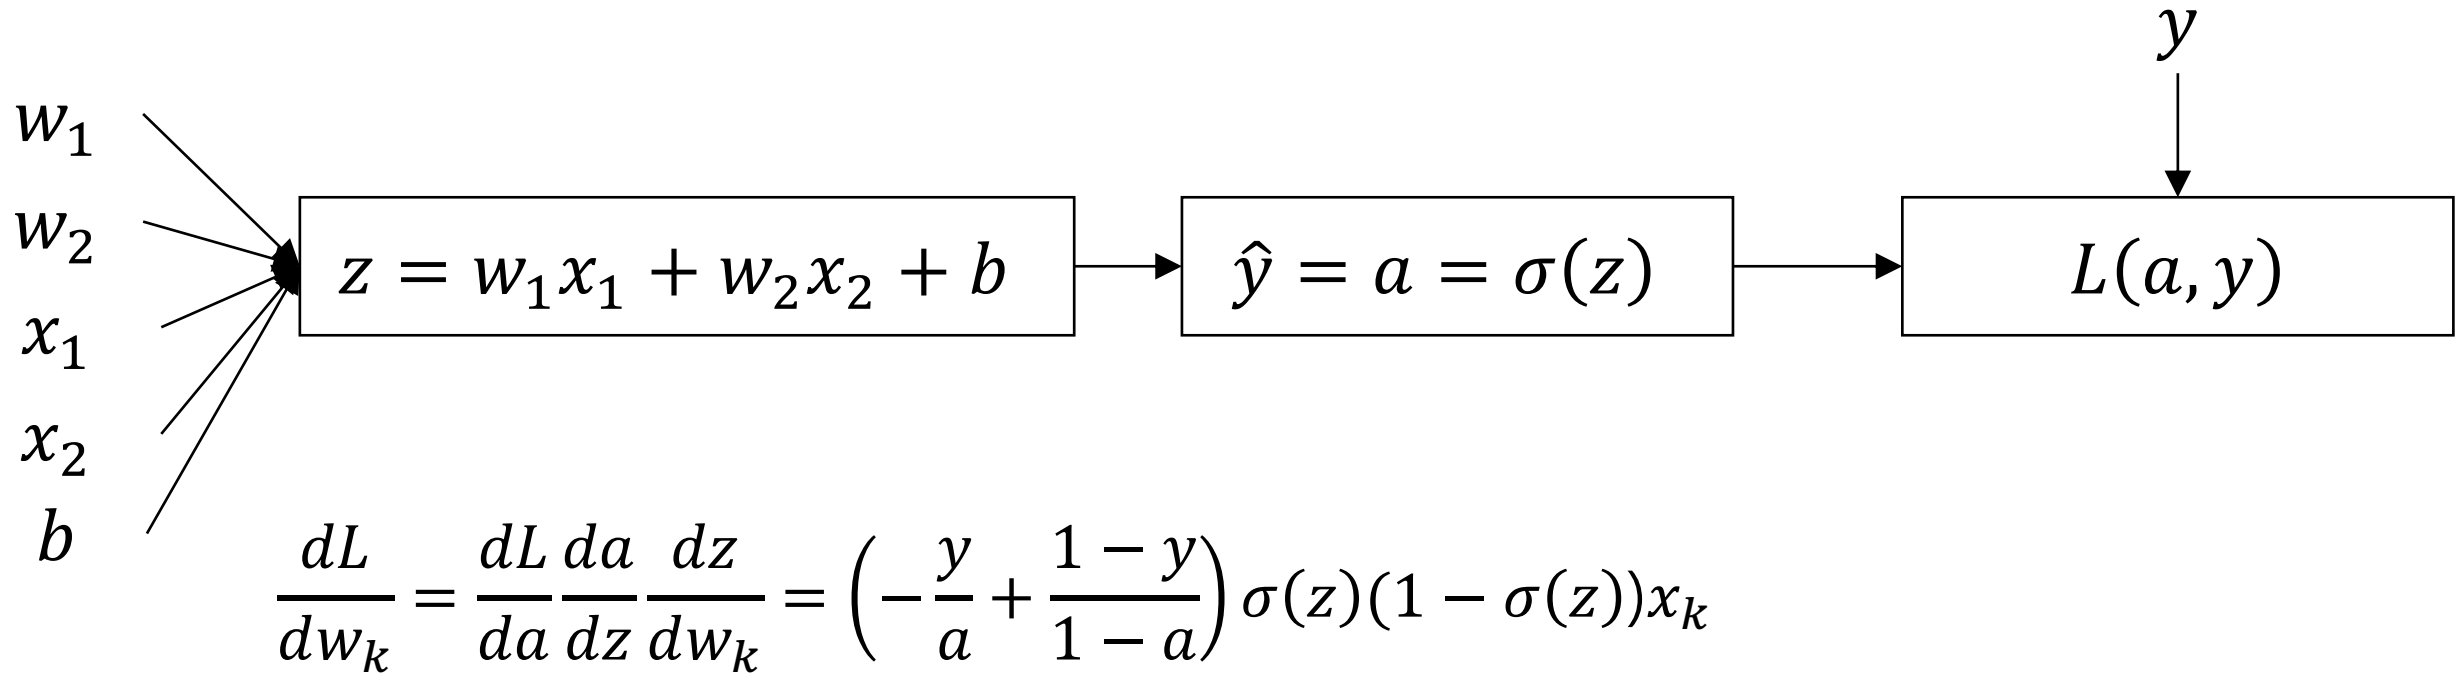
$$L(a, y) = -y \log a - (1 - y) \log(1 - a)$$

Logistic Regression Derivative



$$L(a, y) = -y \log a - (1 - y) \log(1 - a)$$


Logistic Regression Derivative



$$\begin{aligned} \frac{dL}{dw_k} &= \frac{dL}{da} \frac{da}{dz} \frac{dz}{dw_k} = \left(-\frac{y}{a} + \frac{1-y}{1-a} \right) \sigma(z)(1 - \sigma(z))x_k \\ &= \left(-\frac{y}{a} + \frac{1-y}{1-a} \right) a(1-a)x_k = -y(1-a)x_k + (1-y)ax_k \\ &= (a - y)x_k \end{aligned}$$

$$L(a, y) = -y \log a - (1 - y) \log(1 - a)$$

Logistic Regression Derivative



$$\frac{dL}{dw_1} = \frac{dL}{da} \frac{da}{dz} \frac{dz}{dw_1} = (a - y)x_1$$

$$\frac{dL}{dw_2} = \frac{dL}{da} \frac{da}{dz} \frac{dz}{dw_2} = (a - y)x_2$$

$$\frac{dL}{db} = \frac{dL}{da} \frac{da}{dz} \frac{dz}{db} = a - y$$

$$w_1 := w_1 - \eta \frac{dL}{dw_1}$$

$$w_2 := w_2 - \eta \frac{dL}{dw_2}$$

$$b := b - \eta \frac{dL}{db}$$

Gradient descent on m examples


$$J(\mathbf{w}, b) = \frac{1}{m} \sum_{i=1}^m L(\hat{y}^{(i)}, y^{(i)})$$

$$\frac{d}{dw_k} J(\mathbf{w}, b) = \frac{1}{m} \sum_{i=1}^m \frac{d}{dw_k} L(\hat{y}^{(i)}, y^{(i)}) = \frac{1}{m} \sum_{i=1}^m (\hat{y}^{(i)} - y) x_k$$

$$\frac{d}{db} J(\mathbf{w}, b) = \frac{1}{m} \sum_{i=1}^m \frac{d}{db} L(\hat{y}^{(i)}, y^{(i)}) = \frac{1}{m} \sum_{i=1}^m (\hat{y}^{(i)} - y)$$

Logistic Regression on m example

- Initialize \mathbf{w}, b
- $lr = 0.1$
- For $e = 1$ to n_{epoch} :
 - $J = 0$; $d_w1 = 0$; $d_w2 = 0$; $d_b = 0$
 - For $i = 1$ to m :
 - $z = w_1 x_1^{(i)} + w_2 x_2^{(i)} + b$
 - $a = \sigma(z)$
 - $d_w1 += (a - y) x_1^{(i)}$
 - $d_w2 += (a - y) x_2^{(i)}$
 - $d_b += a - y$
 - $w_1 -= lr * d_w1 / m$
 - $w_2 -= lr * d_w2 / m$
 - $b -= lr * d_b / m$



$$\frac{d}{dw_2} J(\mathbf{w}, b) = \frac{1}{m} \sum_{i=1}^m \frac{d}{dw_2} L(a^{(i)}, y^{(i)}) = \frac{1}{m} \sum_{i=1}^m (a^{(i)} - y^{(i)}) x_2$$

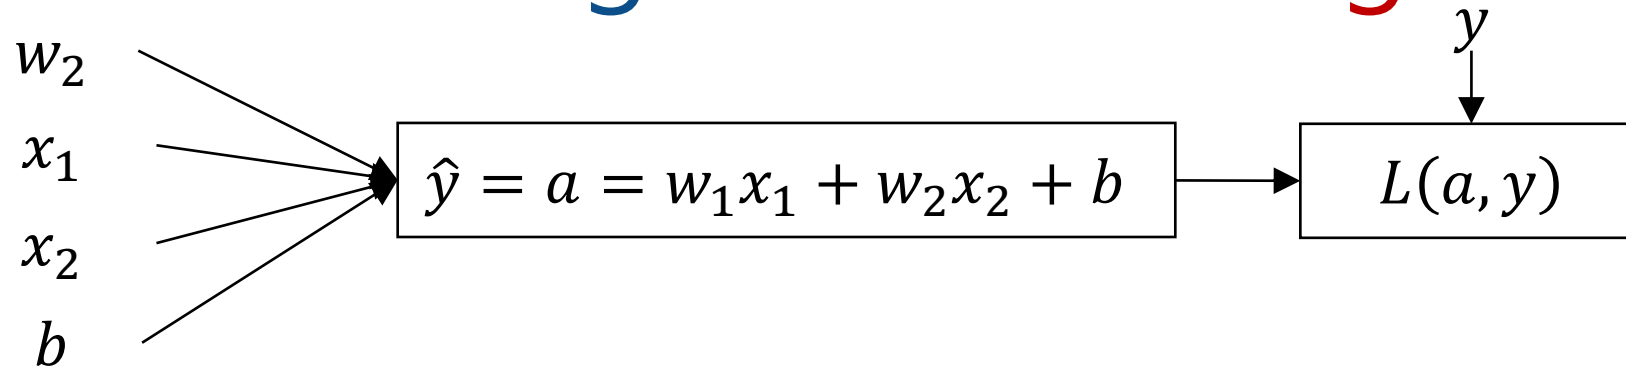
$$\frac{d}{dw_2} J(\mathbf{w}, b) = \frac{1}{m} \sum_{i=1}^m \frac{d}{dw_2} L(a^{(i)}, y^{(i)}) = \frac{1}{m} \sum_{i=1}^m (a^{(i)} - y^{(i)}) x_2$$

$$\frac{d}{db} J(\mathbf{w}, b) = \frac{1}{m} \sum_{i=1}^m \frac{d}{db} L(a^{(i)}, y^{(i)}) = \frac{1}{m} \sum_{i=1}^m (a^{(i)} - y^{(i)})$$

	Linear Regression	Logistic Regression
Problem	Regression	Classification
Model	$\hat{y} = \mathbf{w}^\top \mathbf{x} + b$ Parameters: $\mathbf{w} \in \mathbb{R}^n$, $b \in \mathbb{R}$	$\hat{y} = \sigma(\mathbf{w}^\top \mathbf{x} + b)$ Parameters: $\mathbf{w} \in \mathbb{R}^n$, $b \in \mathbb{R}$
Loss	Squared Error $L(y, \hat{y}) = (y - \hat{y})^2$	Binary Cross Entropy (BCE) $L(y, \hat{y}) = -y \log \hat{y} - (1 - y) \log(1 - \hat{y})$

Cost function: $J(\mathbf{w}, b) = \frac{1}{m} \sum_{i=1}^m L(\hat{y}^{(i)}, y^{(i)})$

Linear regression vs Logistic regression



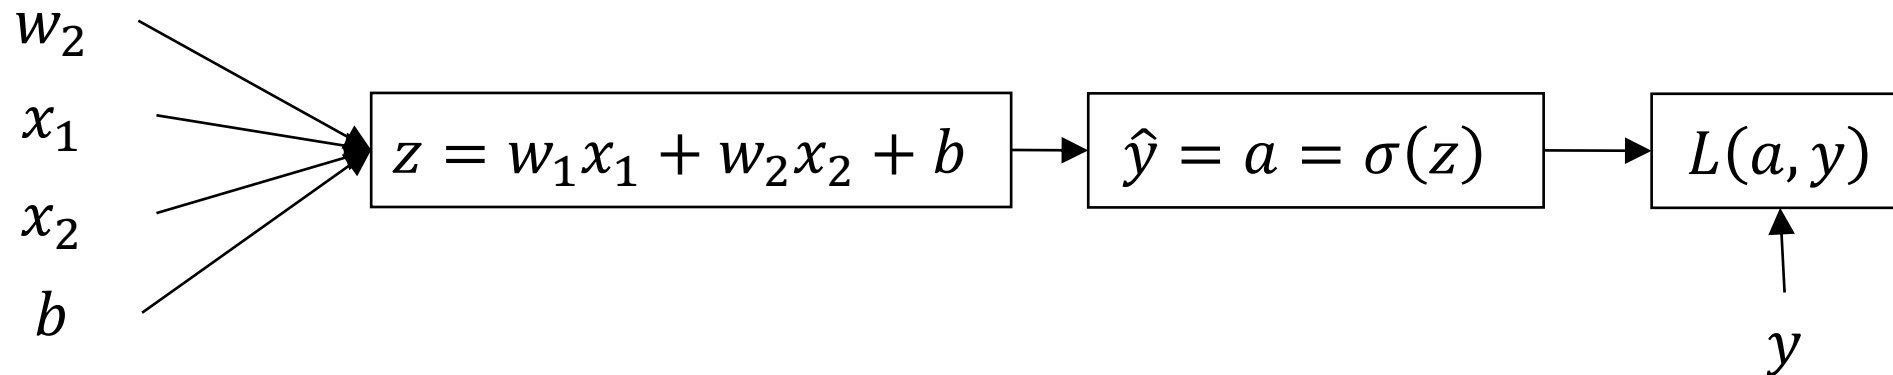
For the simplicity
 $\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$, $\mathbf{w} = \begin{bmatrix} w_1 \\ w_2 \end{bmatrix}$

$$\frac{\partial L}{\partial w_i} = \frac{\partial L}{\partial a} \frac{\partial a}{\partial w_i} = 2(a - y)x_i$$

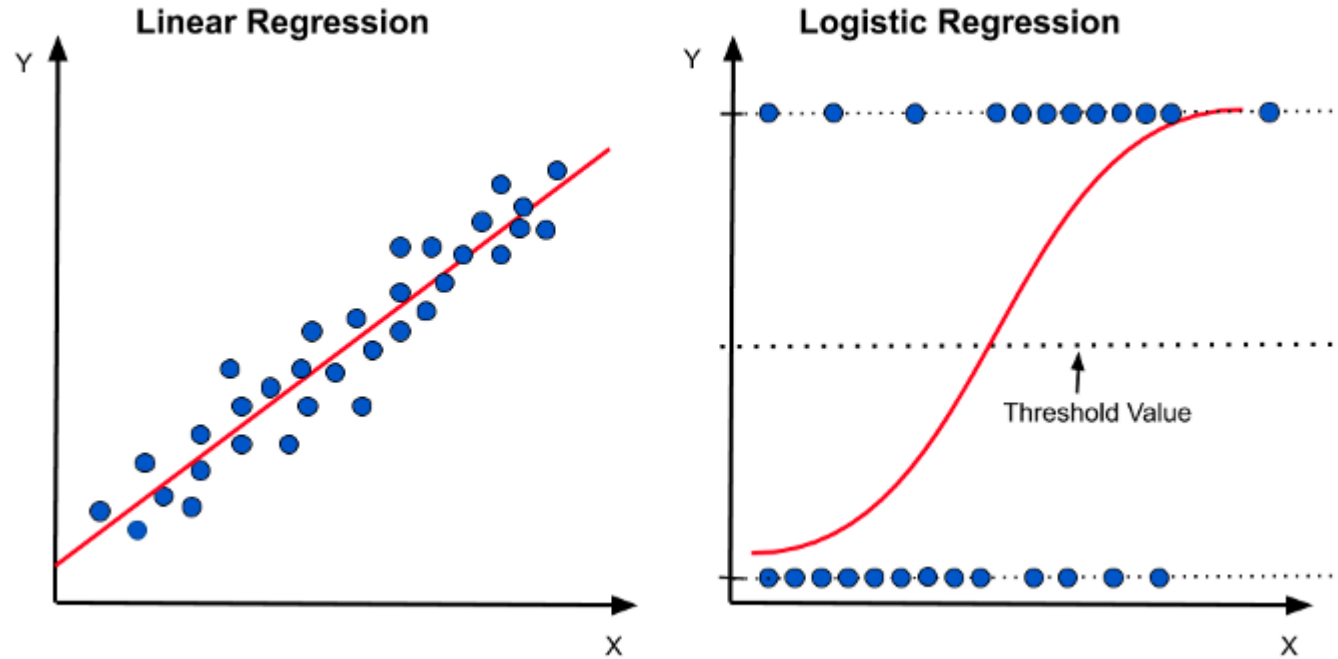
$$\frac{\partial L}{\partial b} = \frac{\partial L}{\partial a} \frac{\partial a}{\partial b} = 2(a - y)$$

$$\frac{\partial L}{\partial w_i} = \frac{\partial L}{\partial a} \frac{\partial a}{\partial z} \frac{\partial z}{\partial w_i} = (a - y)x_i$$

$$\frac{\partial L}{\partial b} = \frac{\partial L}{\partial a} \frac{\partial a}{\partial z} \frac{\partial z}{\partial b} = (a - y)$$



Why is Logistic Regression Called Logistic Regression?



Logistic Regression

Programming in Python

Logistic regression: logical AND

x1	x2	y
0	0	0
0	1	0
1	0	0
1	1	1

$$y = x_1 \text{ AND } x_2$$

Data preparation

Logistic regression (AND)

```
In [ ]: import random  
        from math import exp, log
```

Data prepration

```
In [12]: X = [(0,0),(1,0),(0,1),(1,1)]  
         Y = [0,0,0,1]
```

Model

Model

```
In [14]: class logistic_regression_model():
        def __init__(self):
            self.w = [random.random(), random.random()]
            self.b = random.random()

        def sigmoid(self, z):
            return 1/(1 + exp(-z))

        def predict(self, x):
            z = self.w[0] * x[0] + self.w[1] * x[1] + self.b
            a = self.sigmoid(z)
            return a
```

```
In [15]: model = logistic_regression_model()
```

Training

Training

```
In [16]: def train(X, Y, model, lr = 0.1):
dw0 = 0.0
dw1 = 0.0
db = 0.0
m = len(X)
cost = 0.0
for x,y in zip(X,Y):
    a = model.predict(x)
    if y == 1:
        cost -= log(a)
    else:
        cost -= log(1-a)

    dw0 += (a-y)*x[0]
    dw1 += (a-y)*x[1]
    db += (a-y)

cost /= m
model.w[0] -= lr * dw0/m
model.w[1] -= lr * dw1/m
model.b -= lr*db/m

return cost
```

```
In [17]: for epoch in range(10000):
cost = train(X,Y, model, 0.1)
if epoch % 100 == 0:
    print(epoch, cost)
```

```
0 0.9799277803394626
100 0.4455359447221918
200 0.35278521282410236
300 0.29469845366603453
400 0.25432071172280113
500 0.22425431605184998
600 0.20079558352997323
700 0.1810770070710000
```

...

```
9000 0.019352012397599427
9100 0.019139595049452222
9200 0.018931735521070026
9300 0.018728289777080104
9400 0.018529119755095403
9500 0.01833409306055272
9600 0.018143082680009734
9700 0.01795596671161366
9800 0.017772628111558907
9900 0.017592954455438015
```

Testing

Testing

In [22]: `model.predict((0,0))`

Out [22]: 1.2451625968657186e-05

In [23]: `model.predict((0,1))`

Out [23]: 0.020240526677753723

In [24]: `model.predict((1,0))`

Out [24]: 0.0202405193891944

In [25]: `model.predict((1,1))`

Out [25]: 0.9716510306648906