

Deep Reinforcement Learning : Deep Q-Net (2)



한양대학교 ERICA
소프트웨어융합대학
COLLEGE OF COMPUTING

인공지능학과
Department of
Artificial Intelligence

정 우 환 (whjung@hanyang.ac.kr)

Fall 2022

주요 일정

- 12월 6일 휴강 (서울캠퍼스 출장)
- 12월 15일 기말고사 (수업시간)
- 12월 17일 기말프로젝트 마감

Outline

- PyTorch
 - Save and load models
 - Tensorboard
- 프로젝트 설명
- Review DQN
- Implementing a DQN

Save and Load Models

State_dict

Model's state_dict:

```
conv1.weight      torch.Size([6, 3, 5, 5])
conv1.bias        torch.Size([6])
conv2.weight      torch.Size([16, 6, 5, 5])
conv2.bias        torch.Size([16])
fc1.weight        torch.Size([120, 400])
fc1.bias          torch.Size([120])
fc2.weight        torch.Size([84, 120])
fc2.bias          torch.Size([84])
fc3.weight        torch.Size([10, 84])
fc3.bias          torch.Size([10])
```

모델 정의

```
class TheModelClass(nn.Module):
    def __init__(self):
        super(TheModelClass, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)
```

```
    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 16 * 5 * 5)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```

모델 초기화

```
model = TheModelClass()
```

옵티마이저 초기화

```
optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9)
```

모델의 state_dict 출력

```
print("Model's state_dict:")
for param_tensor in model.state_dict():
    print(param_tensor, "\t", model.state_dict()[param_tensor].size())
```

Save and Load Models

- Save

```
torch.save(model.state_dict(), PATH)
```

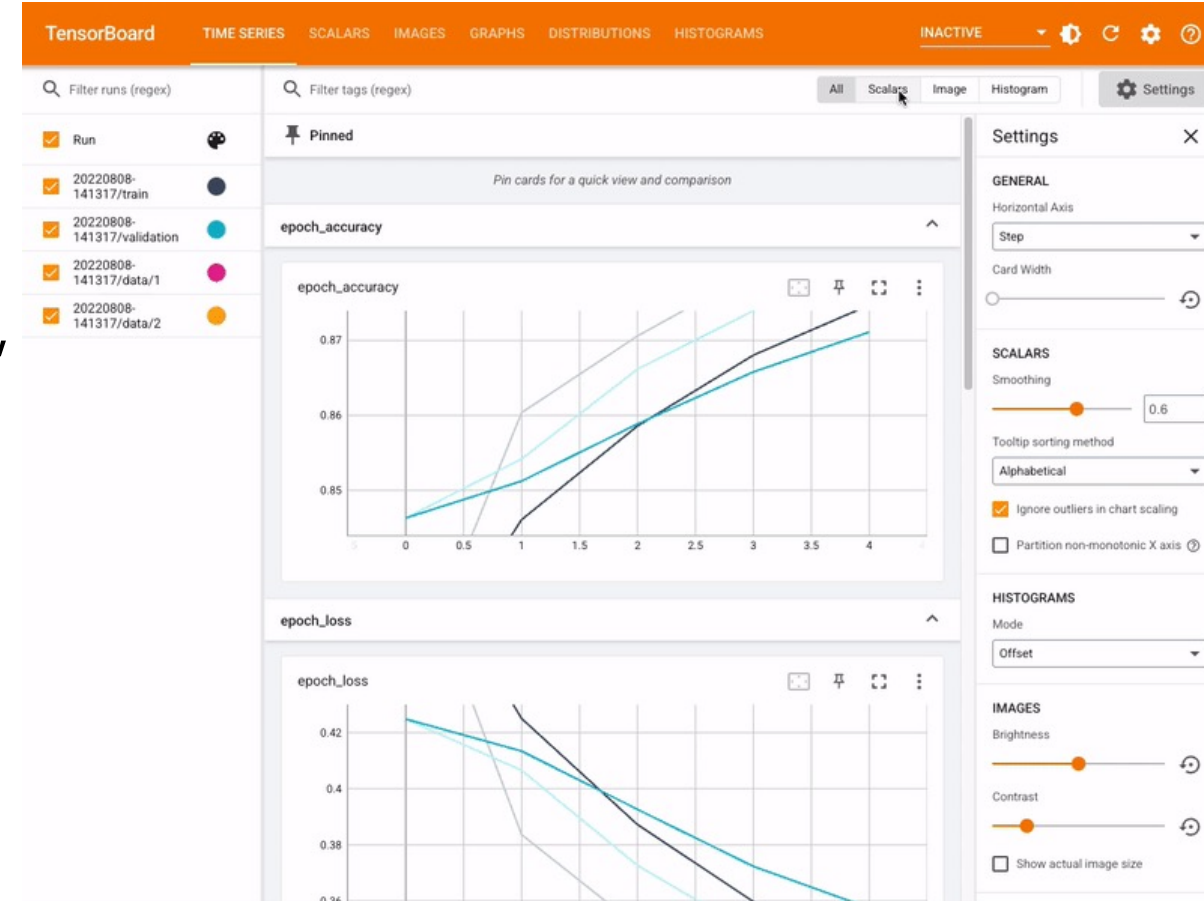
- Load

```
model = TheModelClass(*args, **kwargs)  
model.load_state_dict(torch.load(PATH))  
model.eval()
```

Tensorboard

TensorBoard: TensorFlow의 시각화 툴킷

- Tracking and visualizing metrics such as loss and accuracy
- Visualizing the computation graph
- Viewing histograms of weights, biases, or other tensors as they change over time
- Projecting embeddings to a lower dimensional space
- Displaying images, text, and audio data



Tensorboard

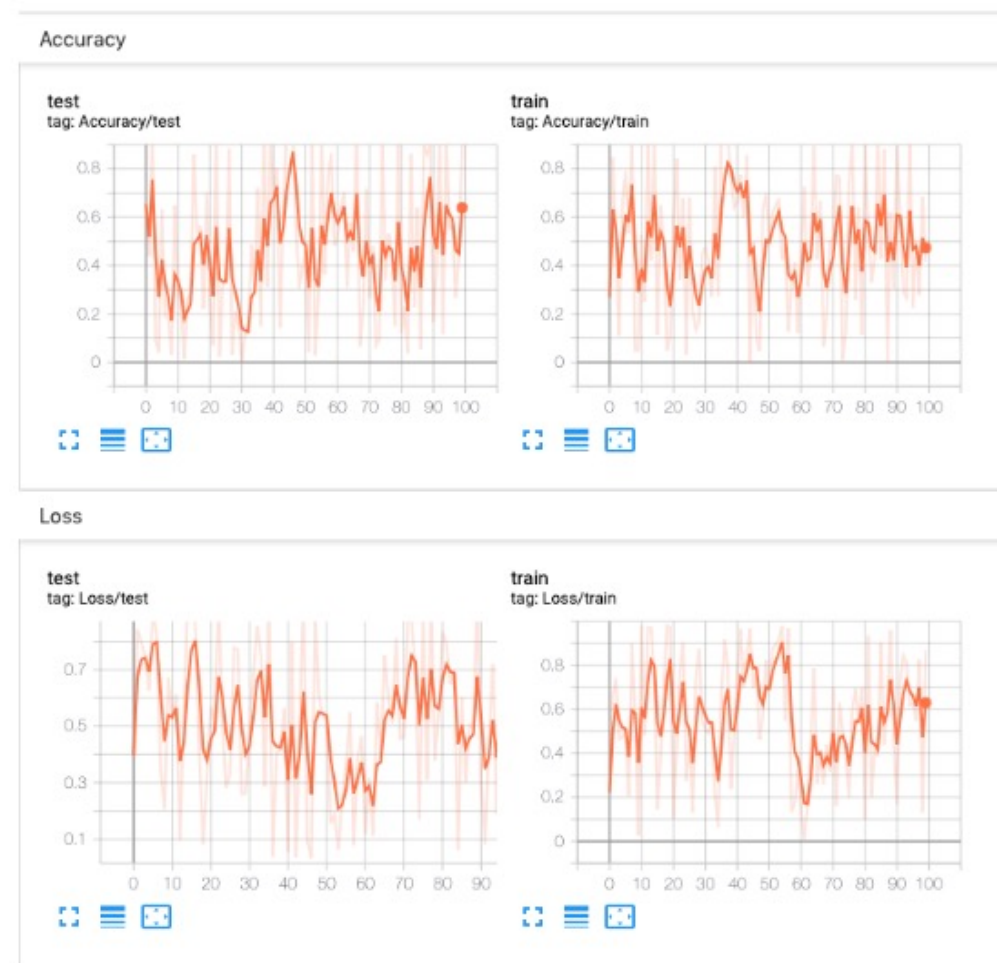
```
pip install tensorboard
tensorboard --logdir=runs
```

```
from torch.utils.tensorboard import SummaryWriter
import numpy as np

writer = SummaryWriter()

for n_iter in range(100):
    writer.add_scalar('Loss/train', np.random.random(), n_iter)
    writer.add_scalar('Loss/test', np.random.random(), n_iter)
    writer.add_scalar('Accuracy/train', np.random.random(), n_iter)
    writer.add_scalar('Accuracy/test', np.random.random(), n_iter)
```

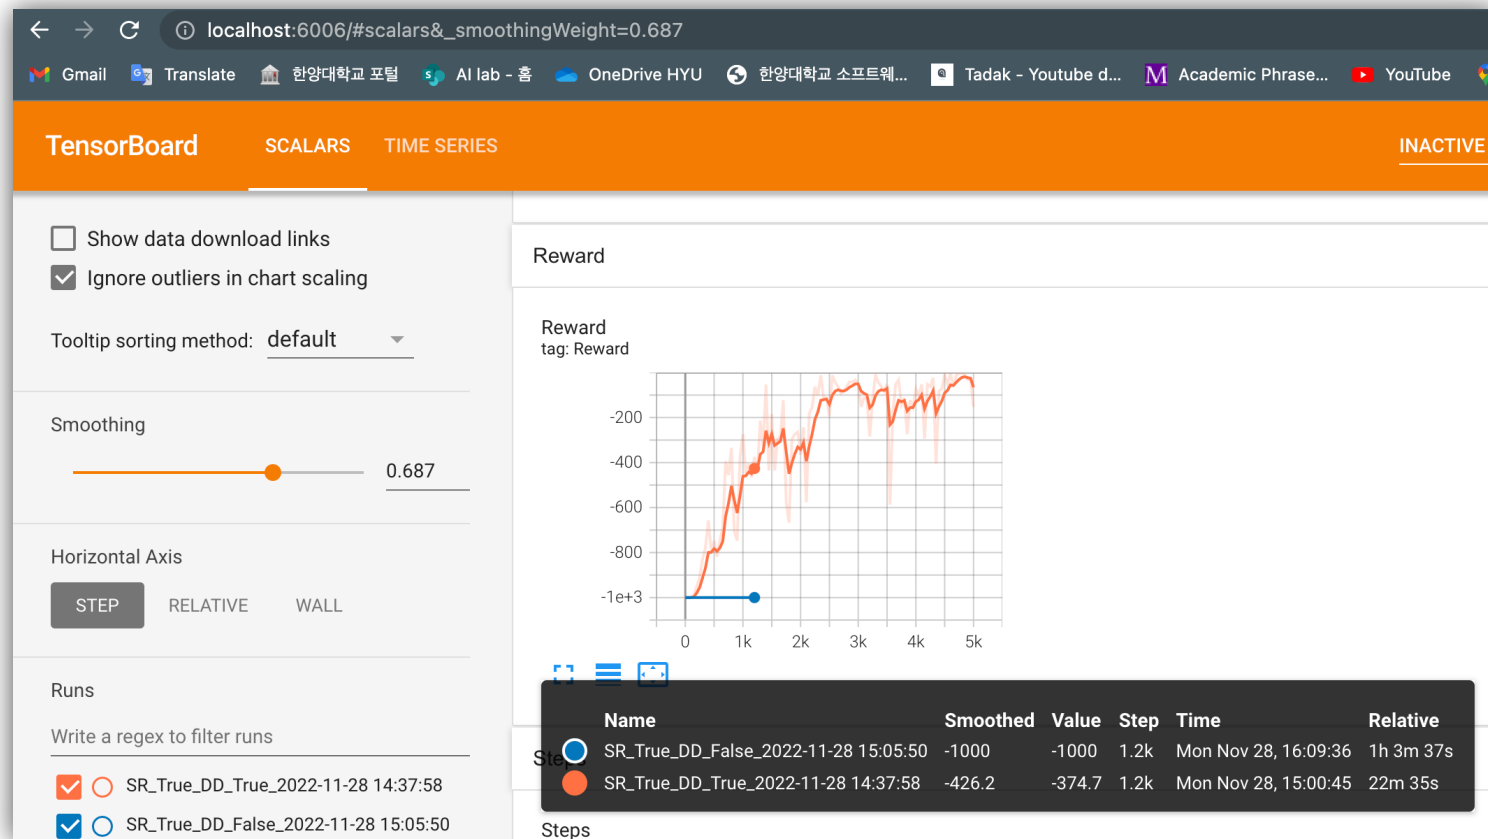
Expected result:



\$ tensorboard --logdir=runs

Tensorboard

```
$ tensorboard --logdir=runs
```



SummaryWriter

```
CLASS torch.utils.tensorboard.writer.SummaryWriter(log_dir=None, comment='', purge_step=None,
max_queue=10, flush_secs=120, filename_suffix='') [SOURCE]
```

Writes entries directly to event files in the `log_dir` to be consumed by TensorBoard.

The *SummaryWriter* class provides a high-level API to create an event file in a given directory and add summaries and events to it. The class updates the file contents asynchronously. This allows a training program to call methods to add data to the file directly from the training loop, without slowing down training.

```
__init__(log_dir=None, comment='', purge_step=None, max_queue=10, flush_secs=120,
filename_suffix='') [SOURCE]
```

Creates a *SummaryWriter* that will write out events and summaries to the event file.

Parameters:

- **log_dir** (*str*) – Save directory location. Default is `runs/CURRENT_DATETIME_HOSTNAME`, which changes after each run. Use hierarchical folder structure to compare between runs easily. e.g. pass in `runs/exp1`, `runs/exp2`, etc. for each new experiment to compare across them.
- **comment** (*str*) – Comment `log_dir` suffix appended to the default `log_dir`. If `log_dir` is assigned, this argument has no effect.
- **purge_step** (*int*) – When logging crashes at step $T + X$ and restarts at step T , any events whose `global_step` larger or equal to T will be purged and hidden from TensorBoard. Note that crashed and resumed experiments should have the same `log_dir`.
- **max_queue** (*int*) – Size of the queue for pending events and summaries before one of the ‘add’ calls forces a flush to disk. Default is ten items.
- **flush_secs** (*int*) – How often, in seconds, to flush the pending events and summaries to disk. Default is every two minutes.
- **filename_suffix** (*str*) – Suffix added to all event filenames in the `log_dir` directory. More details on filename construction in `tensorboard.summary.writer.event_file_writer.EventFileWriter`.

Add scalar data to summary

```
add_scalar(tag, scalar_value, global_step=None, walltime=None, new_style=False, double_precision=False) [SOURCE]
```


Add scalar data to summary.

Parameters:

- **tag** (*str*) – Data identifier
- **scalar_value** (*float* or *string/blobname*) – Value to save
- **global_step** (*int*) – Global step value to record
- **walltime** (*float*) – Optional override default walltime (time.time()) with seconds after epoch of event
- **new_style** (*boolean*) – Whether to use new style (tensor field) or old style (simple_value field).
New style could lead to faster data loading.


```
from torch.utils.tensorboard import SummaryWriter
writer = SummaryWriter()
x = range(100)
for i in x:
    writer.add_scalar('y=2x', i * 2, i)
writer.close()
```

y_2x

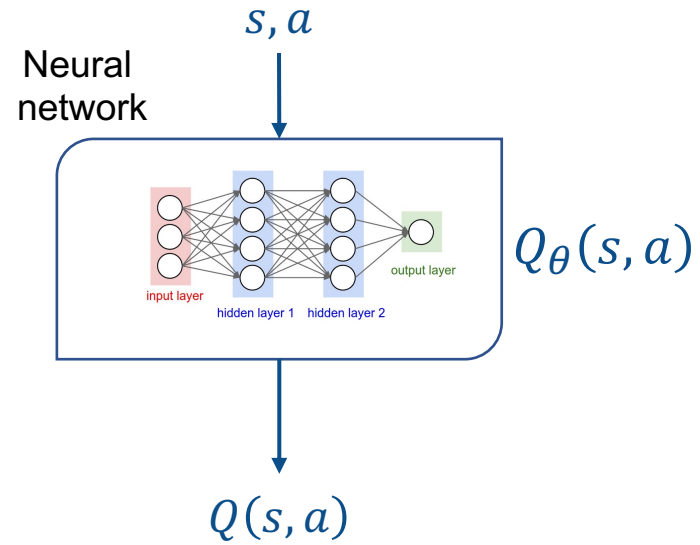


Review

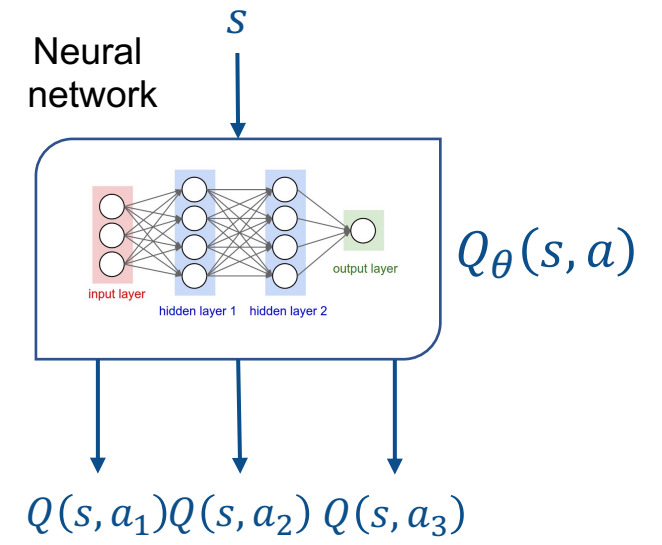
Q-learning (Table vs NN)



Q-learning (Table)



DQN (action-in)



DQN (action-out)

Naïve DQN

Model: $Q_{\theta}(s_t, a_t)$

Training data: $\langle s_t, a_t, r_t, s_{t+1} \rangle$

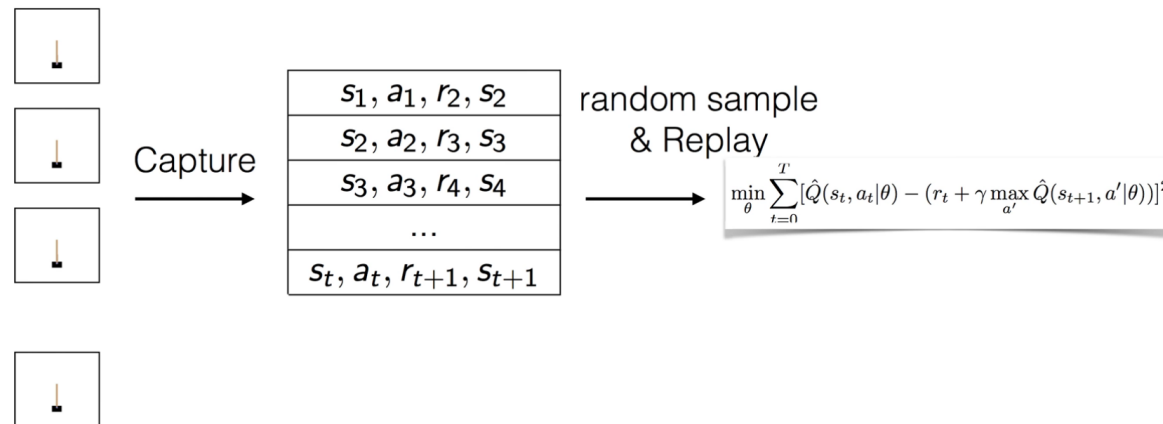
Loss function: $\mathcal{L}(\theta) = \|y_t - Q_{\theta}(s_t, a_t)\|_2^2$ where $y_t = r_t + \gamma Q(s_{t+1}, \pi(s_{t+1}))$

Problems & Solutions

- Naïve Q-learning: Q-table를 Neural Network로 대체
 - Non-stationary target, Correlated samples 문제가 있음

- Solutions

- Capture and replay



- Separate target network

$$\text{minimize } L(\theta) = \left\| \left(r + \gamma \max_{a'} Q_{\theta}(s', a') \right) - Q_{\theta}(s, a) \right\|_2^2$$



$$\text{minimize } L(\theta) = \left\| \left(r + \gamma \max_{a'} Q_{\theta'}(s', a') \right) - Q_{\theta}(s, a) \right\|_2^2$$

Q-learning with experience replay

Initialize replay memory D to capacity N

Initialize action-value function Q with random weights θ

Initialize target action-value function \hat{Q} with weights $\theta^- = \theta$

Initialize the replay memory and two identical Q approximators (DNN). \hat{Q} is our target approximator.

For episode = 1, M **do**

Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequence $\phi_1 = \phi(s_1)$

For $t = 1, T$ **do**

With probability ε select a random action a_t

otherwise select $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$

Execute action a_t in emulator and observe reward r_t and image x_{t+1}

Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in D

Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from D

Set $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$

Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ with respect to the network parameters θ

Every C steps reset $\hat{Q} = Q$

End For

End For

Mnih et al., Human-level control through deep reinforcement learning, Nature 2015 17

Deep Q-net 구현

Taxi

Import

```
1 #!/usr/bin/env python
2 # coding: utf-8
3 import datetime
4
5 import gym
6
7 from time import sleep
8 from tqdm import tqdm
9 import torch
10 from torch import nn
11 import copy
12 import random
13 from torch.utils.tensorboard import SummaryWriter
```

```
24 class Config():...
80
81 class QNet(nn.Module):...
99
100 def train(env, config, qnet):...
124
125 def train_episode(env, config, qnet, qnet2, optimizer, criteria, episode_count = -1):...
168
169 def replay(qnet, qnet2, config, optimizer, criteria, num_instances=5, rec = None):...
202
203 def test(env, config, qnet, global_step = -1):...
248
```

- `class Config()`
 - 학습에 필요한 각종 세팅, replay memory 등을 저장 및 관리
- `class QNet(nn.Module)`
 - DQN class
- `def train(env, config, qnet):`
 - 모델 학습 (전체 에피소드)
- `def train_episode(env, config, qnet, ...):`
 - 모델 학습 (에피소드 1개)
- `def replay(qnet, qnet2, config, ...):`
 - 모델 학습 (에피소드 1개)
- `def test(env, config, qnet,...):`
 - 테스트

Main

```
249 ▶ if __name__ == "__main__":
250     env = gym.make("Taxi-v3").env
251
252     print("Action Space {}".format(env.action_space))
253     print("State Space {}".format(env.observation_space))
254     store_and_replay = True
255     double_dqn = True
256
257     print(f"Store & replay: {store_and_replay}")
258     print(f"Double DQN: {double_dqn}")
259
260     qnet = QNet(env.observation_space.n, env.action_space.n, hidden_dim=32)
261 ● config = Config(store_and_replay=store_and_replay, double_dqn=double_dqn, lr=0.01, gamma=0.75, renew_target=20)
262
263     train(env, config, qnet)
264     test(env, config, qnet, config.num_episodes)
265
```

```

25 class Config():
26     def __init__(self, lr=0.01,
27                 gamma=0.6,
28                 min_epsilon=0.1,
29                 renew_target=100,
30                 store_and_replay=True,
31                 double_dqn=True):
32         self.lr = lr
33         self.gamma = gamma
34         self.min_epsilon = min_epsilon
35         self.store_and_replay = store_and_replay
36         self.double_dqn = double_dqn
37         self.max_time_steps = 1000
38         self.renew_target = renew_target
39
40         self.replay_memory_fail = []
41         self.replay_memory_fail_size = 1000
42         self.replay_memory_success = []
43
44         self.prev_success = set()
45
46         dt = datetime.datetime.now()
47         dt_str = dt.strftime("%Y-%m-%d %H:%M:%S")
48         self.opt_str = f"SR_{store_and_replay}_DD_{double_dqn}"
49         self.writer = SummaryWriter(f"runs/{self.opt_str}_{dt_str}")
50
51         self.num_episodes = 5000

```

```

53 def get_epsilon(self, episode=-1):
54     if episode < 0:
55         return self.min_epsilon
56     return max(self.min_epsilon, 0.99**episode)
57
58 def _insert_record(self, memory, rec, memory_size=-1):
59     if rec in memory:
60         return
61     if len(memory) < memory_size or memory_size == -1:
62         memory.append(rec)
63     else:
64         memory[random.randint(0, memory_size-1)] = rec
65
66 def insert_record(self, rec):
67     done = rec[-1]
68     if done:
69         self._insert_record(self.replay_memory_success, rec)
70         self.prev_success.add(rec[0])
71     else:
72         self._insert_record(self.replay_memory_fail, rec, self.replay_memory_fail_size)
73
74 def get_replay_memory(self):
75     return self.replay_memory_success + self.replay_memory_fail
76
77 def get_replay_record(self):
78     replay_memory = self.get_replay_memory()
79     mid_replay = random.randint(0, len(replay_memory) - 1)
80     return mid_replay, replay_memory[mid_replay]
81

```

class Config():

학습에 필요한 각종 세팅, replay memory 등을 저장 및 관리

class QNet(nn.Module)

DQN class

```
86 class QNet(nn.Module):
87     def __init__(self, num_states, num_actions, hidden_dim = 16):
88         super().__init__()
89
90         self.layers = nn.Sequential(
91             nn.Embedding(num_states, 2*hidden_dim),
92             nn.Linear(2*hidden_dim, 2*hidden_dim),
93             nn.PReLU(),
94             nn.Linear(2*hidden_dim, hidden_dim),
95             nn.PReLU(),
96             nn.Linear(hidden_dim, num_actions)
97         )
98
99     def forward(self, x):
100         # print(x)
101         x = self.layers(x)
102         # print(x)
103         return x
```

def train(env, config, qnet):

모델 학습 (전체 에피소드)

```
105 def train(env, config, qnet):
106     optimizer = torch.optim.SGD(qnet.parameters(), lr=0.001)
107     if config.double_dqn:
108         qnet2 = copy.deepcopy(qnet)
109     else:
110         qnet2 = qnet
111
112     criteria = nn.MSELoss()
113     num_episodes = config.num_episodes
114
115     for i in tqdm(range(1, num_episodes+1)):
116         train_episode(env, config, qnet, qnet2, optimizer, criteria, episode_count=i)
117         if i % config.renew_target == 0:
118             if config.double_dqn:
119                 qnet2 = copy.deepcopy(qnet)
120
121         if i == 1 or i%50 == 0:
122             test(env, config, qnet, i)
123
124
125     print("Training finished.\n")
```

`def train_episode(env, config, qnet, ...):`
모델 학습 (에피소드 1개)

```
while not done and n_steps < config.max_time_steps:
    #n_steps += 1
    state_t = torch.LongTensor([state])

    epsilon = config.get_epsilon(episode_count)
    #print("Epsilon", epsilon)
    if random.uniform(0, 1) < epsilon: # or i < 100:
        action = env.action_space.sample() # Explore action space
    else:
        with torch.no_grad():
            q_hat = qnet(state_t)
            action = torch.argmax(q_hat[0]).item() # Exploit learned values
            # print(q_hat, action)

    next_state, reward, done, info = env.step(action)

    tot_reward += reward

    new_tuple = (state, action, next_state, reward, done)

    state = next_state
    if reward == -10:
        penalties += 1

    # Replay (train)
    if config.store_and_replay:
        config.insert_record(new_tuple)
        loss_i, steps_i = replay(qnet, qnet2, config, optimizer, criteria)
    else:
        loss_i, steps_i = replay(qnet, qnet2, config, optimizer, criteria, rec=new_tuple)
    loss += loss_i
    n_steps += 1
```


`def replay(qnet, qnet2, config, ...):`
모델 학습 (에피소드 1개)

```
177 def replay(qnet, qnet2, config, optimizer, criteria, num_instances=5, verbose=False, rec = None):
178     loss_i = 0
179
180     for _ in range(num_instances):
181         optimizer.zero_grad()
182         if config.store_and_replay:
183             mid_replay, rec = config.get_replay_record()
184
185         state, action, next_state, reward, done = rec
186
187         # Set target value
188         if done:
189             y_t = torch.Tensor([reward])
190         else:
191             next_state_r_t = torch.LongTensor([next_state])
192             with torch.no_grad():
193                 q_next = qnet2(next_state_r_t)
194                 # print("QT", q_target)
195                 y_t = reward + config.gamma * q_next.max(dim=-1)[0]
196
197
198         # Make a prediction
199         state_r_t = torch.LongTensor([state])
200         q_hat = qnet(state_r_t)
201         q_hat = q_hat[:, action]
202
203         # Update
204         loss = criteria(q_hat, y_t)
205         loss.backward()
206         optimizer.step()
207         loss_i += loss.item()
208
209     return loss_i, num_instances
```

def test(env, config, qnet,...):
테스트

```
206 def test(env, config, qnet, global_step=-1):
207     qnet.eval()
208     total_epochs, total_penalties = 0, 0
209     episodes = 100
210
211     total_reward = 0
212     writer = config.writer
213
214     for _ in tqdm(range(episodes)):
215         state = env.reset()
216         epochs, penalties, reward = 0, 0, 0
217
218         done = False
219
220         while not done and epochs < config.max_time_steps:
221             with torch.no_grad():
222                 state_t = torch.LongTensor([state])
223                 q_hat = qnet(state_t)
224                 action = torch.argmax(q_hat[0]).item() # Exploit learned v
225
226                 state, reward, done, info = env.step(action)
227                 total_reward += reward
228                 if reward == -10:
229                     penalties += 1
230
231                 epochs += 1
232
233     total_penalties += penalties
234     total_epochs += epochs
```

```
236     avg_steps = total_epochs / episodes
237     avg_penalty = total_penalties / episodes
238     avg_reward = total_reward / episodes
239
240     print(f"Results after {episodes} episodes:")
241     print(f"Average timesteps per episode: {avg_steps}")
242     print(f"Average penalty per episode: {avg_penalty}")
243     print(f"Average reward per episode: {avg_reward}")
244
245     if global_step > 0:
246         writer.add_scalar("Steps", avg_steps, global_step)
247         writer.add_scalar("Penalty", avg_penalty, global_step)
248         writer.add_scalar("Reward", avg_reward, global_step)
```

```
Run: main ×
▶ ↑ Results after 100 episodes:
■ ↓ Average timesteps per episode: 23.18
⏏ ⌵ Average penalty per episode: 0.0
⏏ ⌶ Average reward per episode: -2.39
⏏ ⌵ |
⏏ ⌶ Process finished with exit code 0
>>
```

References

- Deep Reinforcement Learning, UW CSE Deep Learning – Felix Leeb
- CSCE-689 Reinforcement Learning, Guni Sharon
- Lecture 7: DQN (Sung Kim) <https://youtu.be/S1Y9eys2bdg>
- Mnih et al., Human-level control through deep reinforcement learning, Nature 2015
- Mnih et al., Playing Atari with Deep Reinforcement Learning, Arxiv 2013